



Safe Couplings: Coupled Refinement Types

ELIZAVETA VASILENKO, IMDEA Software Institute, Spain and HSE University, Russia

NIKI VAZOU, IMDEA Software Institute, Spain

GILLES BARTHE, MPI-SP, Germany and IMDEA Software Institute, Spain

We enhance refinement types with mechanisms to reason about relational properties of probabilistic computations. Our mechanisms, which are inspired from probabilistic couplings, are applicable to a rich set of probabilistic properties, including expected sensitivity, which ensures that the distance between outputs of two probabilistic computations can be controlled from the distance between their inputs. We implement our mechanisms in the type system of Liquid Haskell and we use them to formally verify Haskell implementations of two classic machine learning algorithms: Temporal Difference (TD) reinforcement learning and stochastic gradient descent (SGD). We formalize a fragment of our system for discrete distributions and we prove soundness with respect to a set-theoretical semantics.

CCS Concepts: • **Software and its engineering** → **Software verification**; **Functional languages**.

Additional Key Words and Phrases: refinement types, relational types, program verification, Haskell

ACM Reference Format:

Elizaveta Vasilenko, Niki Vazou, and Gilles Barthe. 2022. Safe Couplings: Coupled Refinement Types. *Proc. ACM Program. Lang.* 6, ICFP, Article 112 (August 2022), 29 pages. <https://doi.org/10.1145/3547643>

1 INTRODUCTION

Refinement types provide an appealing mechanism for proving program properties in executable programming languages (including Haskell [Vazou et al. 2014b], Scala [Hamza et al. 2019], and F* [Swamy et al. 2016]). They have been used to good effect for reasoning about functional correctness and termination [Vazou et al. 2014a], resource analysis [Handley et al. 2019], security policies [Lehmann et al. 2021], and other properties of large developments.

However, refinement types do not provide support for reasoning about relational and hyper properties. The main difference between trace properties, which are the usual target of refinement types, and relational and hyper properties is that the latter reason about pairs or sets of traces. This generalization allows to account for a wide variety of security, privacy, and robustness properties.

One natural approach to support relational reasoning is to use relational type systems, as proposed for instance in [Barthe et al. 2014; Maillard et al. 2020]. These type systems are similar to classic refinement type systems, but reason about relational assertions. The latter are interpreted over pairs of (typed) values, and therefore capture relational properties in a natural way. Relational refinement type systems retain the feel of refinement types and are particularly effective when reasoning about two executions of the same program or two programs that follow the same control-flow.

Unfortunately, relational refinement types offer limited support to reason about programs with diverging control-flow. This is due to the fact that relational refinement types are syntax-directed,

Authors' addresses: Elizaveta Vasilenko, elizaveta.vasilenko@imdea.org, IMDEA Software Institute, Madrid, Spain and HSE University, Russia; Niki Vazou, niki.vazou@imdea.org, IMDEA Software Institute, Madrid, Spain; Gilles Barthe, gbarthe@mpi-sp.org, gilles.barthe@imdea.org, MPI-SP, Bochum, Germany and IMDEA Software Institute, Madrid, Spain.



This work is licensed under a Creative Commons Attribution 4.0 International License.

© 2022 Copyright held by the owner/author(s).

2475-1421/2022/8-ART112

<https://doi.org/10.1145/3547643>

whereas many examples of relational program verification benefit from or even require non-syntax-directed reasoning. This is in particular the case for reasoning about program optimizations that restructure the control flow of the program, and about probabilistic programs, since their correctness or security proofs often use mathematical arguments that are not reflected in their syntax. Moreover, it remains a challenge to make relational refinement types practical, even in settings where syntax-directed reasoning suffices. This difficulty is perhaps best witnessed by prior work on relational cost [Çiçek et al. 2019]. In this work, Çiçek et al. [2019] develop BiRelCost, a state-of-the-art bi-directional type checker that compares the cost of two programs or two program executions. In this setting, relational syntax-directed reasoning alternates with non-relational syntax-directed reasoning, in a way that the latter takes over whenever the two program executions no longer have the same control-flow. Unfortunately, controlling such alternations automatically by typing ultimately relies on intricate and partial heuristics. As a consequence, BiRelCost sacrifices predictability and generality, which are some key advantages of refinement types.

A principled approach to overcome the limitations of relational refinement types is to impose a separation between types and relational assertions. This approach, realized by Relational Higher-Order Logic (RHOL) [Aguirre et al. 2017], ensures maximal flexibility and expressiveness. However, the approach is not implemented, thus it remains an open question if RHOL can be made practical.

In this paper, we explore a middle ground approach that retains the key benefits of refinement types. The crux of our approach is a carefully crafted interface for supporting relational reasoning *within* a unary refinement type system. Our approach is implemented atop Liquid Haskell [Vazou et al. 2014b] and inherits many of its essential features: first, our formal guarantees hold for Haskell programs and these programs can be executed using the existing runtime system and optimized libraries of Haskell. Second, verification is carried using a mature refinement type checker and should be familiar for users of Liquid Haskell. Third, the known techniques of Vazou et al. [2018] for encoding proofs manually remain applicable. Naturally, these benefits come at a cost: concretely, our proofs are less automated than proofs in classic Liquid Haskell. However, proofs remain reasonably short, even for relatively complex examples, demonstrating that our middle ground approach achieves predictable and practical verification, a combination that has not been achieved by any prior relational verification tool.

We realize our approach not only for classic higher-order programs, but also for probabilistic programs, an important class of programs that is used pervasively in cryptography, privacy, machine learning, and many other areas. In addition to their numerous applications, probabilistic algorithms are an interesting class of programs to consider in their own right, because they often have intricate specifications and complex proofs. In particular, many properties of interest of probabilistic programs are quantitative, *i.e.* they reason about probabilities or expectations — or in a relational setting, about differences between probabilities or expectations. Although such forms of quantitative reasoning are *seemingly* out of reach of SMT-based verification, prior work has shown that relational verification of probabilistic programs can be achieved using *probabilistic couplings* [Barthe and Hsu 2020; Lindvall 2002; Thorisson 2000; Villani 2009]. We formalize the main tools from coupling-based reasoning in our framework and illustrate how these tools can be used to verify two classic examples of probabilistic programs from machine learning. The first example is Temporal Difference (TD) reinforcement learning, for which we show rapid convergence to a stationary distribution independently of its initial input. The second example is Stochastic Gradient Descent (SGD), for which we show algorithmic stability — a classic machine learning property ensuring that a supervised machine learning algorithm generalizes well and does not overfit with respect to its training set. Both properties are captured in our system as instances of *expected sensitivity*, *i.e.* they upper bound the expected distance between two outputs of the program as a function of the distance between the corresponding inputs. However, both examples use distinct proof tools:

the first example is verified using classic techniques from probabilistic couplings. In contrast, the second example uses a more elaborate, quantitative form of probabilistic couplings which embeds reasoning about the Kantorovich distance [Deng 2015; Villani 2009] between two distributions. Thus, the two examples showcase the different components of our system.

Summary of contributions. Our contributions are the following:

- We define a probabilistic relational refinement type system and encode it into the unary types of Liquid Haskell (§ 3). We choose Liquid Haskell as a mature refinement type checker, but our methodology can be used with other unary refinement type systems.
- We use our system to prove two case studies from the literature: TD (§ 4.1) and SGD (§ 4.2).
- We prove soundness of our type system with respect to a denotational semantics (§ 5).

We start with an overview of our approach (§ 2) and conclude with related work (§ 6).

2 OVERVIEW

We start with an overview of our system that uses (unary) refinement types to machine check relational properties of probabilistic, executable programs. First (§ 2.1) we introduce the `PrM` probabilistic monad and our bins running example. Next, we encode (§ 2.2) and formally prove (§ 2.3) a relational specification for the returned values of bins by axiomatizing probabilistic relational logic as refinement type assumptions. Finally, we follow a similar methodology to encode (§ 2.4) and prove (§ 2.5) a relational property about bins distance.

2.1 The Probability Distribution

The common way to implement probability distributions in Haskell is to use a probability monad, see for instance [Ramsey and Pfeffer 2002]. Therefore, our framework is set up as a verification wrapper around any Haskell library that supports a monadic implementation of probabilities. In order to execute our implementations, we wrapped the `probability` library¹; however, our proofs are independent on the choice of the library and only require the existence of some type `PrM` that implements the standard interface for a probabilistic monad. This includes `pure`, `bind`, and constructors for Bernoulli, choice, and uniform distributions.

`type PrM a = ... -- defined in Sec 3 using an existing probabilistic Haskell library`

```
{-@ type Prob = {p:Double | 0 ≤ p ≤ 1} @-}
{-@ pure      :: a → PrM a @-}
{-@ (>=>)    :: PrM a → (a → PrM b) → PrM b @-}

{-@ bernoulli :: Prob → PrM {v:Integer | v == 0 || v == 1} @-}
{-@ choice   :: Prob → PrM a → PrM a → PrM a @-}
{-@ unif     :: {xs:[a] | 0 < length xs} → PrM a @-}
```

We define the Haskell probability monad `PrM` using an interface of an existing library (§ 3). We use the notation `{-@ ... @-}` to define refinement types and refinement type specifications. That is, the `Prob` type is a Haskell `Double`, refined to be between 0 and 1. The specifications for the monadic `pure` and `>=>` are standard. The `bernoulli` function takes an input a probability `p`, i.e. a `Double` between 0 and 1, and returns 1 with probability `p`, otherwise 0. The function `choice p d1 d2` returns the distribution `d1` with probability `p`, otherwise `d2`. Finally, `unif` takes as input a non-empty list and returns one of its elements uniformly at random. All these functions are

¹We used the probabilistic functional programming library <https://hackage.haskell.org/package/probability-0.2.7>

executed using the underlying Haskell implementation, but are left as uninterpreted (later § 3.4 and § 3.5 axiomatized) in our logic.

Probabilistic Programming: The Bins Example. Using the above interface, we can define (and execute) probabilistic programs. For example below we define the bins program that models a simple balls and bins process. In this process, n balls are thrown into a bin; in each throw, there is a probability that the ball lands outside the bin. The throws are independent and the probability to send any ball in the bin is p . The result of the process is the number of balls that lands into the bin. We model the process using the `bernoulli` distribution.

```
{-@ type Nat = {n:Integer | 0 ≤ n } @-}

{-@ bins :: n:Nat → Prob → PrM {r:Nat | r ≤ n} @-}
bins 0 _ = pure 0
bins n p = do x ← bins (n - 1) p
             y ← bernoulli p
             pure (x + y)
```

We can use standard refinement types to verify various *unary* properties of the `bins` function. In particular, Liquid Haskell will use SMT automation to easily verify that `bins` terminates (because the recursive call occurs on smaller n). One can also prove that the result is always a natural number (as specified by the refined signature) and that it is not greater than n .

Using the theorem proving capabilities of Liquid Haskell [Vazou et al. 2017], we can construct extrinsic proofs that validate probabilistic, unary properties of `bins`. For example, we can define `expect f e` to be the expected value of e , for some function f that turns the values of e to doubles:

```
{-@ expect :: (a → Double) → PrM a → Double @-}
{-@ natToD :: n:Nat → {d:Double | d == to_real n} @-}
```

We can extrinsically prove that `expect natToD (bins n p) = n * p`, assuming that the expectation is linear and `expect natToD (bernoulli p) = p`. Note that the Haskell `Double` is represented, by Liquid Haskell, as `real` in SMT, so the function `natToD` converts natural to double numbers while its specification ensures that the value is not changed. Next, we see how to construct extrinsic proofs that establish relational properties.

2.2 Relational Specifications & Lifting

A first relational property of interest is stochastic dominance, a classic property that defines when a real-valued probabilistic process is better than another. Informally, a real-valued probabilistic process is better than another if it always outputs a “higher value” *w.r.t.* the usual order on real numbers. Interestingly, the intuitive notion of “higher value” is formally defined over two random variables, which makes the definition of stochastic dominance non-trivial. Fortunately, stochastic dominance can be characterized using probabilistic couplings, a classic tool to reason about Markov chains. Informally, couplings are probabilistic equivalent of cartesian products and can be used to lift relational properties to distributions. The lifting of a relational property, formally defined in [Barthe and Hsu 2020], states that two distributions satisfy the lifting of a property p if there exists a coupling of the two distributions such that p holds surely, *i.e.* with probability 1, in this coupling. For our purposes, it suffices to assume an operator \diamond that transforms a relation over two types into a relation over probabilistic distributions over these two types:

```
(◇) :: (a → b → Bool) → PrM a → PrM b → Bool
```

The operator (\diamond) supports compositional reasoning via two axioms: `pureAxiom` states that Dirac distributions of elements related by p are related by the lifted relation $\diamond p$ and `bindAxiom` states that lifted relations are preserved by monadic composition. These axioms, which are expressed below using refinement type signatures, conveniently eschew probabilistic reasoning and open the possibility of carrying SMT-based verification:

```
{-@ assume pureAxiom :: p:(a → b → Bool) → xl:a → xr:b → {p xl xr}
    → {◊ p (pure xl) (pure xr)} @-}

{-@ assume bindAxiom :: p:(b → b → Bool) → q:(a → a → Bool)
    → el:PrM a → fl:(a → PrM b) → er:PrM a → fr:(a → PrM b)
    → {◊ q el er}
    → (xl:a → xr:{a | q xl xr} → {◊ p (fl xl) (fr xr)}
    → {◊ p (el >=> fl) (er >=> fr)} @-}
```

Both `pureAxiom` and `bindAxiom` are defined as Haskell functions that return `()`, but their refinement type signatures encode the desired axioms. The **assume** keyword prevents refinement type checking, since the function definitions do not actually inhabit their type specifications. `pureAxiom` states that for all p, x_l, x_r , if $p \ x_l \ x_r$, then $\diamond p \ (\text{pure } x_l) \ (\text{pure } x_r)$. The type $\{p \ x_l \ x_r\}$ is an abbreviation of $\{v:() \mid p \ x_l \ x_r\}$. In general, we can write $\{q\}$ instead of $\{v:a \mid q\}$, when v does not appear in q . Similarly, the `bindAxiom` ensures $\diamond p \ (e_l \gg= f_l) \ (e_r \gg= f_r)$ when it is provided a proof that $\diamond q \ e_l \ e_r$ and a (higher-order) proof that for all x_l and x_r such that $q \ x_l \ x_r$, $\diamond p \ (f_l \ x_l) \ (f_r \ x_r)$ holds. In § 3 we discuss the implementation of our library that includes these two assumptions, while later (§ 5) we develop a formalisation that justifies these assumptions, concretely, the `pureAxiom` and `bindAxiom` are respectively encoded in the rules T-RET and T-BIND of fig. 8.

In our `bins` example, we want to show that for two throwers sending the same number of balls into bins, the more gifted thrower, *i.e.* the thrower with a higher probability to send balls into the bins, will have a higher count. Formally, our goal is to show that if $p \leq q$, then $\diamond (\leq) \ (\text{bins } n \ p) \ (\text{bins } n \ q)$, from which one can conclude that $\text{expect } \text{naToD} \ (\text{bins } n \ p) \leq \text{expect } \text{naToD} \ (\text{bins } n \ q)$ by a simple property of couplings. In our syntax, we formalize our goal as:

```
{-@ binsSpec :: p:Prob → {q:Prob | p ≤ q} → n:Nat → {◊ (≤) (bins n p) (bins n q)} @-}
```

In the next subsection, we use relational refinement types to establish this goal.

2.3 Relational Proofs

Following Handley et al. [2019], relational proofs are (Haskell) inhabitants of the refinement types that express the relational specifications. Such proofs rely on assumptions about relational properties of the probabilistic primitives. For example, `bins` is using `bernoulli`, thus the proof of `binsSpec` relies on the assumption below, which captures `bernoulli`'s relational specification.

```
{-@ assume bernoulliAxiom :: p:Prob → {q:Prob | p ≤ q}
    → {◊ (≤) (bernoulli p) (bernoulli q)} @-}
```

The specification (formalized in the rule T-BERN of § 5; fig. 8) states that `bernoulli q` stochastically dominates `bernoulli p` if $p \leq q$. In our current implementation, this specification is taken as an axiom, although it would be possible to establish this specification from first principles, by making the definition of lifting explicit for finitely supported distributions.

Using the `bernoulliAxiom` we prove `binsSpec` following the structure of `bins` definition:

```
binsSpec p q 0
  = pureAxiom (≤) 0 0 ()
```

```

binsSpec p q n
= bindAxiom (≤) (≤) (bernoulli p) (bins1 n p) (bernoulli q) (bins1 n q)
  (bernoulliAxiom p q) (\xl xr →
    bindAxiom (≤) (≤) (bins (n-1) p) (bins2 xl) (bins (n-1) q) (bins2 xr)
    (binsSpec p q (n-1)) (\yl yr →
      pureAxiom (≤) (yl + xl) (yr + xr) ())
where
  bins1 n p x = bind (bins (n-1) p) (bins2 x)
  bins2 x y   = pure (y + x)

```

The proof, as `bins`, is inductively defined on n . In the base case, we call `pureAxiom (≤) 0 0 ()` to get $\diamond (\leq) (\text{pure } 0) (\text{pure } 0)$, which concludes the proof, since $\text{bins } 0 p == \text{pure } 0$. Such rewrite steps are automated by Liquid Haskell's logical evaluation strategy (namely PLE [Vazou et al. 2017]). Further, our proofs are automated by SMT arithmetic. For example, `pureAxiom`'s last argument needs to prove that $0 \leq 0$, which is trivially shown by `()` and SMT automation. The inductive case starts with a call to `bindAxiom`, again following `bins` inductive definition. There are two interesting points here. First, the `bins` definition binds `bernoulli p` to a continuation. Since `bindAxiom`'s 4th and 6th arguments require to explicitly pass these continuations, we use a `where` clause to name the continuation `bins1`. Second, `bindAxiom` requires two proof terms. The first one should show that $\diamond (\leq) (\text{bernoulli } p) (\text{bernoulli } q)$, which is shown by the `bernoulliAxiom`. The second one, should show $\diamond (\leq) (\text{bins1 } n p x_l) (\text{bins1 } n q x_r)$, for all x_l and x_r that satisfy $x_l \leq x_r$. We construct such a proof term using a lambda². Since the `bins` definition is using another `bind`, the proof again calls `bindAxiom` with similar arguments. In the last step, `bins` calls `pure`, thus the proof calls `pureAxiom`, whose proof argument is again `()`, i.e. automated by rewriting and SMT.

2.4 Quantitative Specifications and Kantorovich Lifting

So far, we have established that if $p \leq q$ then `bins n q` stochastically dominates `bins n p` and thus $\text{expect natToD } (\text{bins } n p) \leq \text{expect natToD } (\text{bins } n q)$. However, our specification does not provide quantitative information on $\text{expect natToD } (\text{bins } n q) - \text{expect natToD } (\text{bins } n p)$. In fact, one can use simple properties of expectation to show that if $p \leq q$ then we have $\text{expect natToD } (\text{bins } n q) - \text{expect natToD } (\text{bins } n p) \leq n * (q - p)$, where `natToD` is just the cast defined in § 2.1. Unfortunately, one cannot prove this fact using the previous, lifting-based approach. Instead, we need to use a richer notion of lifting that allows to reason about quantitative properties and in particular about the expected distance between two distributions. In order to accommodate such reasoning, one considers a richer setting where each type is equipped with a distance `dist`. These distances are defined inductively on the structure of types; for distribution types, they use the so-called Kantorovich metric [Deng 2015; Villani 2009], which lifts a distance over some types to a distance over its corresponding distribution type:

```

{-@ dist :: Dist a → (a → a → {d:Double | 0 ≤ d}) @-}
{-@ kant :: Dist a → Dist (PrM a) @-}

{-@ kdist :: Dist a → PrM a → PrM a → {d:Double | 0 ≤ d} @-}
kdist d = dist (kant d)

```

The `kant` function turns a distance into Kantorovich and `kdist` simply composes `kant` with `dist`. The formal definition of the Kantorovich metric can be found in Deng [2015]; for this work, it suffices

²The actual proof is using a named function with explicit type specification, since Liquid Haskell does not infer preconditions, but for space, here we use lambdas.

that the Kantorovich metric is based on couplings and lends itself to compositional reasoning. For instance, the following axioms (corresponding to the rules T-RET and T-BIND § 5; fig. 8) are valid:

```
{-@ assume pureDist :: d:Dist a
  → xl:a → xr:a
  → { kdist d (pure xl) (pure xr) == dist d xl xr } @-}

{-@ assume bindDist :: d:Dist b
  → m:Double → p:(a → a → Bool)
  → fl:(a → PrM b) → el:PrM a
  → fr:(a → PrM b) → er::{PrM a | ◇ p el er }
  → { xl:a → {xr:a | p xl xr} → { kdist d (fl xl) (fr xr) ≤ m } }
  → { kdist d (el >>= fl) (er >>= fr) ≤ m } @-}
```

The first axiom states that the Kantorovich distance of two Dirac distributions is the distance of their generating element. The second axiom upper bounds the Kantorovich distance between two monadic compositions by the maximal Kantorovich distance between $f_l x_l$ and $f_r x_r$ for all x_l and x_r related by an intermediate assertion p , such that e_l and e_r are related by the lifting of p . We emphasize that the bind rule for Kantorovich distance is more intricate than the corresponding rules for lifting and that the “obvious” compositional rule that adds distance between the e s and the f s would not be sound, as discussed in § 5.

Then, the distance spec of bins is:

```
{-@ binsDist :: p:Prob → {q:Prob | p ≤ q} → n:Nat
  → { kdist distNat (bins n p) (bins n q) ≤ n * (q - p) } @-}

{-@ distNat :: Dist Nat @-}
```

That is, for each probabilities p and q , so that $p \leq q$, and each natural number n , the Kantorovich distance between $\text{bins } n \ p$ and $\text{bins } n \ q$ is bounded by $n * (q - p)$. Since bins returns natural numbers, the distance is given by distNat that defines the distance metric on natural numbers (§ 3.2). Finally, Haskell’s Doubles are represented in the SMT logic are SMT reals, *i.e.* there is no reasoning about overflows and precision loss. The binsDist specification is well sorted in Z3, since Z3 automatically converts between int and reals.

2.5 Distance Proofs

Unlike the proof of § 2.3, the proof of binsDist is not syntax directed. On the contrary, it requires the construction of a “ghost” probabilistic function that splits the distance between $\text{bins } n \ p$ and $\text{bins } n \ q$. We call this function ghost bins (gbins) and define it as follows:

```
gbins :: Nat → Prob → Prob → PrM Nat
gbins n p q = do x ← bins (n-1) p
               y ← bernoulli q
               pure (x + y)
```

The ghost $\text{gbins } n \ p \ q$ adds bins with probability argument p and bernoulli with probability argument q , thus connecting $\text{bins } n \ p$ and $\text{bins } n \ q$.

Using mostly syntax-directed proofs, we establish the following two lemmata:

```
{-@ binsDistL :: p:Prob → {q:Prob | p ≤ q} → n:Nat
  → { kdist distNat (bins n p) (gbins n p q) ≤ q - p } @-}
```

```
{-@ binsDistR :: p:Prob → {q:Prob | p ≤ q} → n:Nat
    → { kdist distNat (gbins n p q) (bins n q) ≤ (n-1) * (q - p) } @-}
```

Both proofs use the distance axioms of § 2.4. The proof of `binsDistR` (inductively) calls `binsDist`, while the proof of `binsDistL` requires the below axiom for Bernoulli's distance.

```
{-@ assume bernoulliDist :: d:Dist Nat
    → p:Prob → q:Prob
    → { kdist d (bernoulli p) (bernoulli q) ≤ abs (p - q) } @-}
```

That is, the expected distance of two Bernoulli distributions, is bounded by the distance of the `bernoulli`'s arguments (as formalized in the rule T-BERN of § 5; fig. 8).

We prove `binsDist` combining `binsDistL` and `binsDistR` with triangular inequality, which, as explained in § 3.2 is a property (concretely a field) of the `Dist` type. The proof goes by induction:

```
binsDist p q 0
= pureDist distNat 0 0
binsDist p q n
= dist d (bins n p) (bins n q) -- Step 1
  ? trinequality d (bins n p) (gbins n p q) (bins n p) -- Defined in Sec 3.2
=<= dist d (bins n p) (gbins n p q) + dist d (gbins n p q) (bins n q) -- Step 2
  ? binsDistL n p q
  ? binsDistR n p q
=<= (q - p) + (n-1) * (q - p) -- Step 3
=<= n * (q - p) -- Step 4
*** QED
where d = kant distNat
```

The base case is merely an application of the `pureDist` axiom. In the inductive case, we use the (in)equational reasoning proof combinators of [Handley et al. 2019]: `l ? j <= r` ensures `l` is not greater than `r` using the justification `? j` which is optional and `*** QED` concludes the proof. First, we start from the distance between `bins n p` and `bins n q`. Applying triangular inequality, in the second step, we split the distance using `gbins`. Next, we use the two helper lemmata to bound each of the two distances. Finally, using trivial (SMT-automated) arithmetic, we get the desired bound. We note that the lemma `binsDistR` inductively calls `binsDist` on a smaller `n`, so our proof is inductive, while Liquid Haskell is ensuring (mutually recursive) termination.

The `binsDist` example showcases that our framework can be used to machine-check sophisticated proofs that require “ghost” proof objects. To evaluate the expressiveness of our framework, we used it to prove two classic properties of machine learning, probabilistic programs: convergence of TD (§ 4.1) and stability of SGD (§ 4.2).

3 IMPLEMENTATION OF safe-coupling

In this section we present `safe-coupling`³ [Vasilenko and Vazou 2022], a Haskell library that exports an interface for probabilistic programming and permits relational probabilistic verification using Liquid Haskell (§ 3.1). Table 1 summarizes the five main modules of `safe-coupling` that *define* distance (§ 3.2) and the probabilistic monad (§ 3.3), *assume* relational (§ 3.4) and distance (§ 3.5) axioms, and *prove* relational theorems (§ 3.6). In section (§ 5), we formally justify the assumptions made by `safe-coupling`.

³<https://github.com/oquechy/safe-coupling>

Table 1. Summary of safe-coupling library. **LoC** is commented lines of Haskell Code.

	module name	LoC
1.	Data.Dist	155
2.	Monad.PrM	93
3.	Monad.PrM.TCB.Axioms	43
4.	Monad.PrM.TCB.Dist	62
5.	Monad.PrM.Theorems	100
	Total	453

3.1 Liquid Haskell Preliminaries

Verification with Refinement Types. Refinement types are used to do “light” verification. For example `max` of two probabilities (*i.e.* doubles between 0 and 1) is also a probability:

```
max :: Prob → Prob → Prob
max x y = if x ≤ y then y else x
```

To achieve SMT decidable and automatic verification, refinement type systems clearly separate the executable code (here, the Haskell definitions) from the logic (here, the predicates on the refinement types). In the `max` example, every caller of `max` knows the type signature (*i.e.* that the result is also `Prob`), but not its implementation (*i.e.* that it returns one of its arguments). This way verification is modular, but, by default, the Haskell function does not exist in the logical fragment.

User Defined Functions in the Logic. An attempt to refer to user defined definitions in the refinement predicates will lead to an undefined error. For instance, below we define a proof that $\forall x y. x \leq \max x y$ as a function whose arguments encode the quantified x and y and its result is a unit refined with the desired predicate. (Notation: we simplify $\{v : () \mid p\}$ to $\{p\}$.)

```
{-@ not_found :: x:Prob → y:Prob → {x ≤ max x y} @-} -- ERROR: max is unknown
not_found _ _ = ()
```

Since refinement types clearly separate the executable code from the logic, the above specification leads to an error: `max` is unknown to the logic. There are two ways to lift executable definitions in the logic: 1) axiomatization and 2) reflection.

1) *Axiomatization* of `max` defines a logical *uninterpreted* function that has the same refinement type and returns the same result as the executable `max`, but `max`’s definition is not available in the logic. For example, one can use axiomatization to show $0 \leq \max x y \leq 1$ but not that $x \leq \max x y$:

```
{-@ axiomatize max @-}
{-@ ok      :: x:Prob → y:Prob → {0 ≤ max x y} @-} -- max's specification is known
{-@ error   :: x:Prob → y:Prob → {x ≤ max x y} @-} -- max's definition is unknown
```

2) *Reflection* of `max` defines a `max` function in the logic and further makes its definition available:

```
{-@ reflect max @-}
{-@ theorem :: x:Prob → y:Prob → {x ≤ max x y} @-} -- max's definition is known
theorem _ _ = ()
```

Reflection of executable functions permits “deep verification”, *i.e.* reasoning about sophisticated properties like distance of two probabilistic runs. Yet, for decidable refinement type checking, this reasoning requires explicit (user-provided) proofs. Importantly (but not surprisingly) functions can only get reflected, when their definitions consist only of reflected or axiomatized functions. For example, in our setting, functions imported from an unverified Haskell library cannot get reflected.

Proof Terms. Liquid Haskell is using refinement types to encode theorems [Vazou et al. 2018]. The definitions of these functions can be unit (then the theorem is trivially proved by SMT and existing automation) or can contain inductive calls and combine other proof terms using proof combinators. Usually, such definitions do not have runtime meaning: if executed they will not produce any interesting result. But, since Liquid Haskell is checking for totality and completeness of these definitions they encode mathematical proofs.

We call *theorem* a refinement type specification that has a proof term, like the theorem for `max` defined above. As an alternative, Liquid Haskell's `assume` keyword lets you assume *axioms* (refinement types) that one cannot prove. We use such axioms to encode properties of axiomatized functions. For example, when `max` is axiomatized we can define two axioms that describe its behavior.

```
{-@ axiomatize max @-}
{-@ assume max1 :: x:Prob → y:Prob → {x ≤ max x y} @-}
{-@ assume max2 :: x:Prob → y:Prob → {y ≤ max x y} @-}
```

Since `max` is axiomatized, its definition is not available in the logic, so none of the above axioms can be proved.

3.2 Data.Dist: Definition of Distance

We used Liquid Haskell to define the refined data type `Dist a` that encodes a metric, as follows.

```
{-@ data Dist a = Dist {
    dist      :: a → a → {v:Double | 0.0 ≤ v}
  , identity  :: x:a → {dist x x == 0.0}
  , symmetry  :: x:a → y:a → {dist x y == dist y x}
  , trinequality :: x:a → y:a → z:a → {dist x z ≤ dist x y + dist y z}
} @-}
```

The first field of `Dist` contains the distance function `dist` on any expressions of type `a`: it is a function that given two arguments of type `a` returns a non negative `Double`. The next three fields capture the metric's axioms for identity of indiscernibles, symmetry, and triangle inequality.

In this module, we further defined the distance metric on doubles and natural numbers:

```
distDouble :: Dist Double
distNat     :: Dist Nat
```

These definitions contain both the definition of the function `dist` and the proofs of the metric axioms on the concrete distance. Further, we defined a function that computes distance between two same-length lists of a given a `Dist a`:

```
{-@ dList :: Dist a → xs:[a] → ys:[a] | length xs == length ys → {v:Double | 0 ≤ v} @-}
```

We proved all the metric axioms of `dList`, yet, since there exists the same-length dependency it is not possible to define a (well-typed) `Dist [a]` function. Still, we use the above definitions to compute distance between natural numbers, doubles, and their lists:

```
assert (dist distDouble 42.0 40.0 == 2.0)
assert (dList distDouble [42] [40.0] == 2.0)
```

3.3 Monad.PrM: Definition of the Probabilistic Monad

The module `Monad.PrM` is essentially a wrapper around an executable Haskell probability monad. We chose the probabilistic functional library `probability`, due to its clear interface. Our development uses `probability` to execute (and test) our probabilistic programs, but our mechanized proofs do not depend on it and could use alternative libraries (e.g. `monad-bayes` [Scibior et al. 2015]).

The underlying probability library (here prefixed as `PLib`) exports the type `T prop a`, that essentially maps each `a` to a probability `prop` and defines the monadic and probabilistic primitives.

The data type. Our probability monad type instantiates `prop` to the probability type `Prob`.

```
type PrM a = PLib.T Prop a
```

Axiomatized primitives. For each monadic (`>=>` and `pure`) and probabilistic primitive (`bernoulli`, `uniform`, and `choice`) operations we used the probability functions to define the Haskell function and axiomatized (as described in § 3.1) them in the logic. For example, `pure` is defined as follows:

```
{-@ axiomatize pure :: x:a → PrM {v:a | v == x} @-}
pure x = PLib.pure x
```

We followed this encoding for practical reasons: since `PLib` is not itself verified with Liquid Haskell (which is a challenging future work) its definitions are not available in the logic. Yet, this encoding leaves us the flexibility to axiomatize the primitives as desired (§ 3.4 and 3.5).

Reflected functions. Using the axiomatized primitives we defined probabilistic functions, e.g. `mapM`.

```
{-@ reflect mapM :: (a → PrM b) → [a] → PrM [b] @-}
```

Since `mapM` is reflected, its definition (which is standard) is available in the logic and used to prove (relational) theorems about `mapM` (§ 3.6).

3.4 TCB.Axioms: Assumption of Relational Axioms

The first trusted computing base (TCB) of safe-coupling is called `Axioms` and contains the relational specification of each axiomatized primitive. It provides the Haskell axiomatized function (`◇`) and uses it to encode the relational axioms for `pure`, (`>=>`), and `bernoulli`, as presented in § 2.2.

3.5 TCB.Dist: Assumption of Distance Specifications

The second trusted computing base of safe-coupling is called `Dist` and contains the distance specification of each axiomatized primitive. It provides the Haskell function `kant` that (like `◇`) is axiomatized in the logic and for each distance `d` on `a` returns the Kantorovich distance on distributions of `a` and (as defined in § 2.4) `kdist` that simply composes `kant` with `dist`:

```
{-@ axiomatize kant :: Dist a → Dist (PrM a) @-}
{-@ kdist :: Dist a → PrM a → PrM a → {d:Double | 0 ≤ d} @-}
```

This module provides the distance axioms for the axiomatized primitives. In § 2.4 we presented the axiomatization for `bind` (`bindDist`), `pure` (`pureDist`), and `bernoulli` (`bernoulliDist`). We assume three more axioms presented in fig. 1. First, `unifDist`, the distance axiom of `unif`, states that the Kantorovich distance between two uniform distributions is zero, when the sampling input lists are equal up to a permutation. Second, `choiceDist`, the distance axiom for `choice`, states that the Kantorovich distance of two choice expressions `choice p el ul` and `choice q er ur` is `p` times the Kantorovich distance of `el` `er` and `1 - p` times the Kantorovich distance of `ul` `ur`, when `p = q`. Finally, `pureBindDist` is a distance axiom for `bind`. For soundness reasons discussed in § 5, the rule is stated only for `bind` expressions whose second argument is (the monadic lifting of) a pure function. The axiom requires that the pure functions `fl` and `fr` make the distance between two values grow by at most `m`; in other words, the distance between `fl xl` and `fr xr` cannot exceed the distance between `xl` and `xr` and some fixed constant `m`. Under this assumption, the Kantorovich distance between `(el >=> (ppure . fl))` and `(er >=> (ppure . fr))` is bounded by the Kantorovich distance between `el` and `er` plus `m`. This specialized axiom provide a means to upper bound the Kantorovich distance

```

{-@ assume unifDist :: d:Dist a → xsl::{[a] | 1 ≤ length xsl}
    → xsr::{[a] | isPermutation xsl xsr}
    → { kdist d (unif xsl) (unif xsr) == 0 } @-}

{-@ assume choiceDist :: d:Dist a → p:Prob
    → el:PrM a → ul:PrM a → er:PrM a → ur:PrM a
    → { kdist d (choice p el ul) (choice p er ur) ≤
        p * (kdist d el er) + (1.0 - p) * (kdist d ul ur) } @-}

{-@ assume pureBindDist :: da:Dist a → db:Dist b → m:Double
    → fl:(a → b) → el:PrM a → fr:(a → b) → er:PrM a
    → (xl:a → xr:a → {dist db (fl xl) (fr xr) - dist da xl xr ≤ m})
    → { kdist db (el >>= (ppure . fl)) (er >>= (ppure . fr)) ≤
        kdist da el er + m } @-}

```

Fig. 1. Distance Axioms of safe-coupling. Encoding rules T-UNIF, T-CHOICE, and T-BINDRET of fig. 8.

between two bind expressions as a function of the Kantorovich distance of their first arguments, and is instrumental for our case studies.

3.6 Theorems: Proof of Relational Properties

This module proves common theorems using our assumed TCB and the defined functions of `Monad.PrM`. Concretely, it provides a simplified version of the `bindAxiom` when the two bind arguments form a bijectional coupling and a relational specification for the monadic map.

All the properties on this module are proved. Later (§ 5), we provide a formalism that justifies the assumptions of our two TCB modules.

4 CASE STUDIES

To evaluate safe-coupling we used it to verify two classic machine learning properties: convergence of TD (§ 4.1) and stability of SGD (§ 4.2). § 4.3 summarizes our results.

4.1 Case Study I: Convergence of TD(0)

Our first case study proves convergence for TD(0), a classical algorithm for Reinforcement Learning.

4.1.1 Implementation of TD(0). In the standard reinforcement learning setting, an agent (*i.e.* the learning algorithm) repeatedly interacts with the environment, a Markov Decision Process (MDP) with state space `State` and set of actions `A`. At each step, the MDP reacts to the agent's action by drawing a new random state and a numeric reward according to a function $t :: \text{State} \rightarrow \text{PrM}(\text{State}, \text{Double})$. The current state i of the process is known to the learner, but the exact function t is not. Given black-box access to t , the goal of the learner is to find a policy map $\pi :: \text{State} \rightarrow A$ from the state space to the best available action from `A` that maximizes the learner's expected reward over infinite time.

Figure 2 presents the implementation of TD(0), a Temporal Difference (TD) learning algorithm that estimates the value function $v :: \text{State} \rightarrow \text{Reward}$ of the MDP, *i.e.* the expected reward at each state if the agent were to repeatedly act according to some assumed policy π . For simplicity of verification, we defined `State` as $\{s:\text{Nat} \mid s \leq n\}$ and functions on `State` as lists of length n . The TD learner (*i.e.* `td0`) takes as input the number of iterations n , a transition function t , and an

```

import Monad.PrM -- mapM defined here

td0 :: Int → [PrM (State, Reward)] → [Reward] → PrM [Reward]
td0 n t v = iterate n (act t) v

act :: [PrM (State, Reward)] → [Reward] → PrM [Reward]
act t v = mapM (sample t v) [0..length v]

sample :: [PrM (State, Reward)] → [Reward] → State → PrM Reward
sample t v i = do (j, r) ← t !! i
                  pure (update v i j r)

iterate :: Int → (a → PrM a) → a → PrM a
iterate 0 _ x = pure x
iterate n f x = f x >>= iterate (n - 1) f

update :: [Reward] → State → State → Reward → Reward
update v i j r = (1 -  $\alpha$ ) * (v !! i) +  $\alpha$  * (r +  $\gamma$  * v !! j)

type State = Int
type Reward = Double

```

Fig. 2. Implementation of TD(0).

estimate of v and it iterates through the n states. At each iteration i , the learner runs `sample` that draws a reward and transition (j, r) from the i th element of t . Then, the estimate v_j is updated by incorporating the observed reward r and the estimated value v_j of the new state. Estimated rewards in the future are reduced by the factor $\gamma \in [0, 1)$. Higher γ allows v to converge faster.

4.1.2 Convergence for TD(0). Our goal is to show that `td0` converges to a stationary distribution independently of its initial input. This can be achieved by proving that `td0` is contractive on v . One potential approach would be to prove that for $k = \alpha * \gamma + (1 - \alpha)$,

$$\{-@ \text{td0Goal} :: n:\text{Nat} \rightarrow t:[\text{PrM (State, Reward)}] \rightarrow v_l:[\text{Reward}] \rightarrow v_r:[\text{Reward}] \rightarrow \{k\text{dist dList (td0 n } v_l \text{ t) (td0 n } v_r \text{ t)} \leq k^n * (\text{dist dList } v_l \text{ } v_r)\} @-\}$$

Instead, we prove a stronger property that `td0` is contractive for all possible outcomes (via lifting). To do so, first we defined a pure (to be lifted) predicate that bounds the distance (since Liquid Haskell does not permit lambdas in the refinements):

```

bounded :: Dist a → Double → a → a → Bool
bounded d m v1 v2 = dist d v1 v2 ≤ m

```

Using `bounded` we define the `td0` specification as follows:

$$\{-@ \text{td0Spec} :: n:\text{Nat} \rightarrow t:[\text{PrM (State, Reward)}] \rightarrow v_l:[\text{Reward}] \rightarrow v_r:[\text{Reward}] \rightarrow \{\diamond (\text{bounded dList (} k^n * (\text{dist dList } v_l \text{ } v_r)) \text{ (td0 n } v_l \text{ t) (td0 n } v_r \text{ t)}) @-\}$$

The `td0Spec` specification implies our original `td0Goal`. This is because a bound property on all outcomes implies an average bound:

$$\{-@ \text{bound} :: d:\text{Dist a} \rightarrow k:\text{Double} \rightarrow e_l:\text{PrM a} \rightarrow e_r:\text{PrM a}$$

```
type PDouble = {Double | 0 ≤ m}
```

```
{-@ iterateSpec :: m:PDouble → n:Nat → f:([Reward] → PrM [Reward])
  → (m:PDouble → xl: [Reward] → xr: [Reward] →
    {bounded dList m xl xr ⇒ ◇ (bounded dList (m * k)) (f xl) (f xr)})
  → rl: [Reward] → rr: [Reward]
  → {bounded dList m rl rr ⇒ ◇ (bounded dList (m * kn)
    (iterate n f rl) (iterate n f rr))} @-}

{-@ actSpec :: m:PDouble → t:[PrM (State, Reward)] → vl: [Reward] → vr: [Reward]
  → {bounded dList m vl vr ⇒ ◇ (bounded dList (k * m)) (act t vl) (act t vr)} @-}

{-@ sampleSpec :: m:PDouble → t:[PrM (State, Reward)] → vl: [Reward] → vr: [Reward]
  → i:State
  → {bounded dList m vl vr ⇒ ◇ (bounded distDouble (k * m))
    (sample t vl i) (sample t vr i)} @-}

{-@ updateSpec :: vl: [Reward] → vr: [Reward] → i:State → j:State → r:Reward
  → {distD (update vl i j r) (update vr i j r) ≤
    k * max (distD (vl !! i) (vr !! i)) (distD (vl !! j) (vr !! j))} @-}
```

Fig. 3. Relational Lemmas for the `td0Spec` Proof; where $\text{distD } x \ y = \text{dist distDouble } x \ y$.

$$\rightarrow \{ \diamond (\backslash x_l \ x_r \rightarrow \text{dist } d \ x_l \ x_r \leq k) \ e_l \ e_r \Rightarrow k \text{dist } d \ e_l \ e_r \leq k \} @-\}$$

The reverse implication does not hold: two distributions can have a Kantorovich distance that is upper bounded by k and do not satisfy the lifting of bounded $d \ k$. As a counterexample, assume $e_l = [(3, 0.5), (77, 0.5)]$ and $e_r = [(7, 0.5), (75, 0.5)]$, then for $k = 3$ and d the distance of natural numbers, the right side is true (i.e. $k \text{dist distNat } e_l \ e_r \leq 3$), but the left side does not hold. Thus, and to our surprise, we could prove a stronger property than originally anticipated, and in a simpler system with plain, non-quantitative liftings.

4.1.3 Proof of Convergence for TD(0). We proved `td0Spec` in 128 lines of (Liquid) Haskell code and, as summarized in table 2, we used five lemmas; one for each used function. We named each lemma by postfixing the name of the function with `Spec`. The specification `mapMSpec` comes from the `safe-coupling` library (§ 3.6), while the remaining lemmas are presented in fig. 3.

The proof of each lemma, like `binsSpec` of § 2.3, is following the structure of the function definition. Concretely, `td0Spec` is proved by `iterateSpec`, using `actSpec` as the proof requirement; `iterateSpec` is proved by induction, using the pure and bind axioms; `actSpec` is proved by `mapMSpec`, using `sampleSpec` as the proof argument; `sampleSpec` is using the axioms and `updateSpec`, which is proved using the linearity and triangular inequality of the distance on doubles.

In short, the great challenge was to come up with the correct invariant for `td0Spec`, after which the proof follows the structure of the `td0`'s implementation.

4.2 Case Study II: Stability of SGD

Supervised machine learning algorithms are algorithms that aim to select the best fitting model from a class of parametric models by iteratively refining some initial parameter, based on some training set. A good measure of the quality of these algorithms is their, so called, generalization


```

import Monad.PrM
import Data.Derivative (grad)

{-@ sgd :: zs:[DataPoint] | 2 ≤ length zs → w0:Weight → αs:[StepSize]
    → f:(DataPoint → Weight → Double) → PrM Weight @-}
sgd _ w0 [] _ = ppure w0
sgd zs w0 (α:αs) f = do z ← unif zs
                        w ← sgd zs w0 αs f
                        pure (update z α f w)

update :: DataPoint → StepSize → (DataPoint → Weight → Double)
      → Weight → Weight
update z α f w = w - α * grad (f z) w

{-@ type StepSize = {v:Double | 0 ≤ v} @-}
type StepSize     = Double
type DataPoint    = (Double, Double)
type Weight       = Double

```

Fig. 4. Implementation of SGD.

error, which measures how they perform on previously unseen data. One sufficient condition for an algorithm to have a controlled generalization error is to be algorithmically stable [Bousquet and Elisseeff 2002]. Informally, a supervised machine learning algorithm is algorithmically stable if the output of the algorithm is not overly dependent on any single element in the training set. More formally, a supervised machine learning algorithm is ϵ -stable if the Kantorovich distance between the parameters obtained by running the algorithm on the same initial parameter and two adjacent training sets (*i.e.* training sets that differ in a single element) is upper bounded by ϵ . In this case study, we show that Stochastic Gradient Descent, the *de facto* backpropagation algorithm for deep learning, is algorithmically stable. The proof follows the steps of [Hardt et al. 2016].

4.2.1 Implementation of SGD. Figure 4 presents our `sgd` implementation, a variant of Stochastic Gradient Descent. The algorithm takes as input a training set `zs`, modeled as a list of data points, an initial weight `w0`, a list of learning step sizes `αs`, and loss function `f`. In the general setting, the loss function `f` in the definition of SGD would be a vector function $Z \times \mathbb{R}^d \rightarrow \mathbb{R}$, where Z is the data set and \mathbb{R}^d is the weight space. For the sake of simplicity, in our implementation `Weight` is defined as a single value of type `Double` which corresponds to $d = 1$. We assume that `f` is L-Lipschitz.

The function `sgd` recursively computes a sequence of so-called weights (or classifiers) starting from the initial parameter `w0`, by updating at each step the current weight `w` into `update z α f w`, where `z` is sampled uniformly from data set `zs` and the learning step `α` represents the influence of each iteration in the final result.

In our proof, we unfold the definition of uniform sampling based on the operator choice. The unfolding is possible when `length zs` is at least 2, which is encoded as a precondition in `sgd`.

```

{-@ unif :: {xs:[a] | 0 < length xs} → PrM a @-}
unif :: [a] → PrM a
unif [a]      = pure a
unif zs@(x : xs) = choice (1.0 / natToD (length zs)) (pure x) (unif xs)

```

```

{-@ assume contractive :: d:Dist Double → α:StepSize
    → f:(DataPoint → Weight → Double) → z:DataPoint
    → wl:Weight → wr:Weight
    → {dist d (update z α f wl) (update z α f wr)
       == dist d w w'} @-}

{-@ assume bounded :: d:Dist Double → α:StepSize
    → f:(DataPoint → Weight → Double)
    → zl:DataPoint → zr:DataPoint → wl:Weight → wr:Weight
    → {dist d (update zl α f wl) (update zr α f wr)
       == dist d wl wr + 2 * lip * α} @-}

```

Fig. 5. Assumptions of the `sgdDist` Proof.

To implement update, we assume a partial function `grad :: (Weight → Double) → Weight → Double` which computes the gradient of `f`. Our proof does not rely on `grad`'s definition, therefore it can be imported from any automatic differentiation library.

4.2.2 Stability of SGD. The stability statement bounds Kantorovich distance of two runs of `sgd` on data sets `zsl` and `zsr`, which differ in exactly one element. We encode this by requiring the existence of data points `x` and `y` and a common tail `zs` such that `(x:zs)` and `(y:zs)` are permutations of the given datasets. We express the stability statement as a refinement type specification, as follows.

```

{-@ sgdDist :: d:Dist Double
    → x:DataPoint → y:DataPoint → zs:[DataPoint] | 1 ≤ length zs
    → zsl::{DataSet|isPermutation (x:zs) zsl} → wsl:Weight
    → zsr::{DataSet|isPermutation (y:zs) zsr} → wsr:Weight
    → αs:StepSizes → f:LossFunction
    → {dist (kant d) (sgd zsl wsl αs f) (sgd zsr wsr αs f) ≤
       dist d wsl wsr + estab (length zsl) αs} @-}

estab :: Nat → [StepSize] → Double
estab l αs = 2.0 * lip / natToD l * sum αs

```

That is, the Kantorovich distance between two runs of `sgd` is bounded by the distance of the different weights plus an ϵ , defined by the helper function `estab`. Note that `estab` does not depend on the different inputs of `sgd`, but depends on `lip`, which represents the Lipschitz constant L of function `f` and is axiomatized in our proof.

4.2.3 Assumptions. Proofs of algorithmic stability are traditionally based on strong assumptions on the loss function $f : Z \times \mathbb{R}^d \rightarrow \mathbb{R}$, namely:

- (1) *L-Lipschitz*: $\|f(z, w_1) - f(z, w_2)\| \leq L\|w_1 - w_2\|$;
- (2) *Convex*: $f(z, w_1) \geq f(z, w_2) + \langle \nabla f(z, w_2), w_1 - w_2 \rangle$; and
- (3) *β -smooth*: $\|\nabla f(z, w_1) - \nabla f(z, w_2)\| \leq \beta\|w_1 - w_2\|$ and for all step sizes $\alpha, \alpha\beta < 1$.

where $\|\cdot\|$ and $\langle \cdot \rangle$ are the norm and the scalar product on \mathbb{R}^d respectively.

Instead of directly encoding these assumptions, which require advanced mathematical machinery that is not readily available in Liquid Haskell, we assume two properties of the update function. These properties follow from the assumptions on the loss function and are presented in fig. 5.

```

{-@ assume leftId :: x:a → f:(a → PrM b) → { (pure x) >>= f == f x } @-}

{-@ assume assoc  :: x:PrM a → g:(a → PrM b) → f:(b → PrM c)
    → { (x >>= g) >>= f == x >>= (\x → g x >>= f) } @-}

```

Fig. 6. Monad laws.

The contractiveness axiom for update (*contractive*) states that the distance of update on the same data point is equal to the distance of the weights. The boundedness of the difference of gradients (*bounded*) states that the distance of update on different data points is equal to the distance of the weights plus $2 * \text{lip} * \alpha$.

Lastly, we assume two *monad laws* (fig. 6). There is a convention that every Haskell monad adheres to monad laws, two of which we use in the proof of stability. Left identity law states that for every $x:a$ and probabilistic function $f:(a \rightarrow \text{PrM } b)$, the distribution $f \ x$ is equal to a Dirac distribution of x sequentially composed with f . The associativity law allows us to compose two sequential probabilistic computations f and g into a single computation $\lambda x \rightarrow g \ x \gg= f$. In our proof, after the unfolding of *unif*, these two laws once again alter the program structure.

4.2.4 Proof of Stability of SGD. Using the assumptions described above, we proved *sgdDist* in 251 lines of Liquid Haskell code. The proof proceeds by structural induction on αs . Further, we applied three program transformations to the *@sgd@* definitions.

First, in the inductive case of *sgd*, sampling from *unif* zs_l is replaced with *unif* $(x:zs)$ by using *unifDist* axiom and identity property of distance.

Second, *sgd* $zs_l \ ws_l \ \alpha s \ f$ is further transformed by unfolding the definition of *unif*. For the next modification, we use the property of distributivity of choice over *bind*:

```

{-@ choiceBind :: p:Prob → e_l:PrM a → e_r:PrM a → f:(a → PrM b)
    → {choice p e_l e_r >>= f = choice p (e_l >>= f) (e_r >>= f)} @-}

```

This property is proved by applying the *assoc* law and unfolding the definition of *choice*:

```

choice :: Prob → PrM a → PrM a → PrM a
choice p a b = bernoulli p >>= \x → if x == 1 then a else b

```

After rewriting with *choiceBind*, the inductive case of *sgd* is brought into the following shape:

```

sgd zs w0 ( $\alpha:s$ ) f
= choice (1.0 / natToD (length (x:zs)))
    (pure x >>= (\z → sgd (x:zs) w0  $\alpha s$  f >>= (pure . update z  $\alpha$  f)))
    (unif zs >>= (\z → sgd (x:zs) w0  $\alpha s$  f >>= (pure . update z  $\alpha$  f)))

```

Lastly, the *leftId* law is applied to simplify the first branch of *choice*.

We use the above structure of *@sgd@* to distinguish between two cases that correspond to the two branches of *choice*. In the first possibility, where the two algorithms sample the same element from zs , we apply the contractiveness property of update. Whereas, in the case where the two algorithms sample the elements in which the datasets differ (x and y), we apply boundedness of the difference of gradients. The *choiceDist* axioms then guarantees that Kantorovich distance increases by $2 * \text{lip} * \alpha / n$ at each iteration, from which we conclude by induction. Other than *choiceDist*, our proof is using the *pureDist* and *pureBindDist* axioms, respectively in the base and inductive case.

Table 2. Summary of case studies. **LoC** is lines of executable Haskell code to define the implementations. **Proof LoC** is lines of commented Liquid Haskell code to define refinement type specifications and their proofs. **Lemmas** is the number of total lemmas proved. **Axioms** is the assumed specifications. **Time** is verification time in seconds on a 2,8 GHz Dual-Core Intel Core i5, RAM 8 GB 1600 MHz DDR3.

Case study	LoC	Proof LoC	Lemmas	Axioms	Time (sec)
bins Spec (§ 2.3)	21	28	3	0	8
bins Dist (§ 2.5)		143	7	0	113
td0 (§ 4.1)	31	128	5	0	24
sgd (§ 4.2)	22	251	7	3	51
Total	74	550	22	3	196

4.3 Quantitative Summary

Table 2 summarizes the effort required to verify the three examples that we presented through the paper: bins, td0, and sgd. In total, we used 550 lines of proof code, *i.e.* commented (Liquid) Haskell lines that express refinement types specifications and their proofs, to verify 74 lines of executable Haskell code, giving an executable-to proof-ratio of almost 6, which is high but expected, given that our proofs are extrinsic. Most of our proofs directly follow the structure of the definitions, thus are easy, once the proper specification is set. The exception is the distance proof for bins (§ 2.5) which required the construction of a “ghost” proof distribution, providing evidence that such sophisticated proofs are feasible in our framework. The same proof is an outlier for our verification time: it requires almost 2 minutes, while the rest of the proofs need less than 1 minute. We note that for td0 we distributed the proof over multiple Haskell modules to allow fast and interactive proof development, since Liquid Haskell verifies per-module and provides local proof-error messages.

5 PROOF SYSTEM

To justify the axioms of our implementation, in this section, we define λ^{RP} , a core probabilistic λ -calculus, with a set of relational proof rules. Although each proof rule of the relational program logic is encoded independently from others as an axiom in Liquid Haskell, we follow the same style of presentation as Aguirre et al. [2017] and treat our set of proof rules as a proof system. This treatment is primarily motivated by our desire to prove a crisp statement for soundness. We also note that for the purpose of establishing soundness, we only consider discrete distributions.

This section is organized as follows: we define the syntax (§ 5.1) and type system (§ 5.2) of λ^{RP} ; then, the axioms (§ 5.3) and the proof rules (§ 5.4) of our logic; and finally, the denotational semantics of λ^{RP} (§ 5.5) which we use to show soundness (§ 5.6).

5.1 Syntax

We consider a typed probabilistic λ -calculus with algebraic datatypes and distributions (fig. 7). Types are built from base types using the usual function space constructor and type constructors for lists (`list`) and probability distributions (`prM`; which encodes the Haskell type `PrM` of § 3.3).

λ^{RP} features a rich set of *constants* that include natural (n) and real (a) numbers, the special constant $+\infty$, the true and false booleans, the `nil` and `cons` list constructors. Furthermore, λ^{RP} features constants for arithmetic operations, boolean operations, equality, and inequality.

Variables in λ^{RP} include three special variables d , r_l , and r_r that respectively model distance and the left and right result of computations.

Types	$t, s ::= \text{nat} \mid \text{bool} \mid \text{list } t$	<i>Discrete Types</i>
	$\mid \text{real} \mid t \rightarrow t \mid \text{prM } t$	<i>Continuous Types</i>
Constants	$c ::= \text{nil} \mid \text{cons} \mid n \in \mathbb{N} \mid a \in \mathbb{R}^+ \mid +\infty \mid \text{true} \mid \text{false}$	
	$\mid + \mid - \mid \cdot \mid / \mid = \mid < \mid \wedge \mid \vee \mid \Rightarrow \mid \neg$	
Expressions	$e, u ::= \text{unif } e \mid \text{bern } x \mid \text{choice } x \ e \ e$	<i>Probabilistic</i>
	$\mid \text{ret } e \mid \text{bind } e \ e$	<i>Monadic</i>
	$\mid \lambda x. e \mid e \ e \mid c \mid x$	<i>Pure</i>
	$\mid \text{case } e \text{ of } \{\text{nil} \rightarrow e; \text{cons } x \ x \rightarrow e\}$	
	$\mid \text{let } x = e \text{ in } e$	
Assertions	$p ::= e \mid p[\wedge, \vee, \Rightarrow] p \mid \neg p$	<i>Boolean</i>
	$\mid \forall x : t. p \mid \exists x : t. p$	<i>Quantifiers</i>
	$\mid \diamond_e p \ e \ e$	<i>Lifting</i>
Environments		
Typing	$\Gamma ::= \emptyset \mid x:t, \Gamma$	
Predicate	$\Phi ::= \emptyset \mid p, \Phi$	

Fig. 7. Syntax of λ^{RP} . (Variables include r_l , r_r , and d .)

Expressions are built from constants and variables using the standard λ -calculus and monadic constructions. The former include constructions for lambda abstraction ($\lambda x. e$), application ($e \ e$), case analysis ($\text{case } e \text{ of } \{\text{nil} \rightarrow e; \text{cons } x \ x \rightarrow e\}$), and structurally recursive definitions ($\text{let } x = t \text{ in } e$). The latter include the probabilistic primitives $\text{unif } e$ that probabilistically returns an element of its input list e , with uniform distribution, $\text{bern } x$ that returns 1 with probability x , otherwise 0, and $\text{choice } x \ e_1 \ e_2$ that returns the distribution e_1 with probability x and e_2 with probability $1 - x$. These primitives model our implementation interface (§ 3.3).

Assertions include (arbitrary, but boolean typed) expressions of λ^{RP} and boolean operators. To allow reduction to RHOL [Aguirre et al. 2017], assertions also include quantifiers even though our system does not explicitly use them. Finally, λ^{RP} has the lift assertion $\diamond_k p \ e_l \ e_r$ that encodes the combination of (\diamond) and kdist of our implementation (§ 3). Here k is a real typed expression, p is a relational assertion on two arguments of type t_l and t_l respectively, and e_l and e_r are probability distributions over t_l and t_l respectively. The assertion ensures that p holds for the distribution coupling and that the Kantorovich distance between the two distributions is bounded by k .

We adopt standard conventions, e.g. $g \cdot f$ stands for $\lambda x. g \ (f \ x)$. By abuse of notation, we write $\text{ret } \cdot f$ as shorthand for $\lambda x. \text{ret } (f \ x)$ and $\diamond_k p$ as syntactic sugar for $\diamond_k (\lambda r_l \ r_r. p) \ r_l \ r_r$, i.e. when the lifting happens on the special variables r_l and r_r . By convention, we also write $\diamond p$ as a shorthand for $\diamond_{+\infty} p$. As usual, we also let $e[e_x/x]$ denote the capture-free substitution of e_x for x in e .

5.2 Type System

We equip our language with a simple type system which serves three purposes: first, it ensures that expressions respect the type signatures of operators; second, it ensures that recursive definitions are structurally terminating. Our logic is agnostic to the mechanism used to enforce structural termination, so we leave this mechanism abstract. Finally, our type system restricts the use of distribution types to discrete types, so that types and expressions of our language can be given a set-theoretic interpretation—we discuss the case of continuous distributions in the § 7.

5.3 Axioms

The special variable d encodes distance that we assume satisfies the axioms of an (extended) metric:

Relational Probabilistic Typing

$$\begin{array}{c}
\boxed{\Gamma; \Phi \vdash e_l : \text{prM } t_l \sim e_r : \text{prM } t_r \mid \diamond_k p} \\
\hline
\frac{}{\Gamma; \Phi \vdash e_l : \text{prM } t_l \sim e_r : \text{prM } t_r \mid \diamond_{+\infty} \text{true}} \text{T-BOT} \\
\frac{r_l : t_l, r_r : t_r, \Gamma; \Phi \vdash p \Rightarrow p_w \quad r_l : t_l, r_r : t_r, \Gamma; \Phi \vdash k \leq k_w}{\Gamma; \Phi \vdash e_l : \text{prM } t_l \sim e_r : \text{prM } t_r \mid \diamond_k p} \text{T-WEAKEN} \\
\frac{\Gamma; \Phi \vdash e_l : \text{prM } t_l \sim e_r : \text{prM } t_r \mid \diamond_{k_w} p_w}{\Gamma; \Phi \vdash \text{unif } e_l : \text{prM } t \sim \text{unif } e_r : \text{prM } t \mid \diamond_0 r_l = r_r} \text{T-UNIF} \\
\frac{\Gamma; \Phi \vdash x_l : \text{real} \sim x_r : \text{real} \mid 0 \leq r_l \leq r_r \leq 1}{\Gamma; \Phi \vdash \text{bern } x_l : \text{prM nat} \sim \text{bern } x_r : \text{prM nat} \mid \diamond_{|x_r - x_l|} 0 \leq r_l \leq r_r \leq 1} \text{T-BERN} \\
\frac{\Gamma; \Phi \vdash x_l : \text{real} \sim x_r : \text{real} \mid 0 \leq r_l = r_r \leq 1}{\Gamma; \Phi \vdash e_l : \text{prM } t_l \sim e_r : \text{prM } t_r \mid \diamond_{k_e} p \quad \Gamma; \Phi \vdash u_l : \text{prM } t_l \sim u_r : \text{prM } t_r \mid \diamond_{k_u} p} \text{T-CHOICE} \\
\frac{\Gamma; \Phi \vdash \text{choice } x_l e_l u_l : \text{prM } t_l \sim \text{choice } x_r e_r u_r : \text{prM } t_r \mid \diamond_{x_l \cdot k_e + (1-x_l) \cdot k_u} p}{\Gamma; \Phi \vdash e_l : t \sim e_r : t \mid p \ r_l \ r_r \wedge d(r_l, r_r) \leq k} \text{T-RET} \\
\frac{\Gamma; \Phi \vdash \text{ret } e_l : \text{prM } t \sim \text{ret } e_r : \text{prM } t \mid \diamond_k p}{\Gamma; \Phi \vdash e_l : \text{prM } s \sim e_r : \text{prM } s \mid \diamond_k q} \\
\frac{x_r : s, x_l : s, \Gamma; q \ x_l \ x_r, \Phi \vdash f_l \ x_l : \text{prM } t \sim f_r \ x_r : \text{prM } t \mid \diamond_{a \cdot d(x_l, x_r) + b} p}{\Gamma; \Phi \vdash \text{bind } e_l f_l : \text{prM } t \sim \text{bind } e_r f_r : \text{prM } t \mid \diamond_{a \cdot k + b} p} \text{T-BIND}
\end{array}$$

Fig. 8. Typing of λ^{RP} , where k ranges over distance (real typed expressions). We use the syntactic sugar $g \cdot f \doteq \lambda x. g(f\ x)$ and $\diamond_k p \doteq \diamond_k(\lambda r_l \ r_r. p) \ r_l \ r_r$.

Definition 5.1 (Metric Axioms). For every type t and every x, y, z of type t :

- (1) *Identity:* $d(x, y) = 0 \Leftrightarrow x = y$.
- (2) *Symmetry:* $d(x, y) = d(y, x)$.
- (3) *Triangular Inequality:* $d(x, z) \leq d(x, y) + d(y, z)$.

We assume that the distance d is defined for all types of λ^{RP} . For nat and real it is defined as $d(x, y) = |x - y|$. For booleans it is defined as $d(x, y) = 1$ if $x \neq y$, i.e. it coincides with the discrete distance on booleans. For lists of equal length, we assume the distance is the maximum of distances between elements at the same positions. When the length is different, the distance is infinite. For functions, it is defined as the maximum distance over all function domains and for distributions, as the Kantorovich distance (of Villani [2009]). These assumptions are required for soundness.

5.4 Proof System

Our proof system uses two judgments to decide logical implication and relational typing. The first judgment $\Gamma; \Phi \vdash p$ states that the assertion p is valid under the assumptions of Φ , where all the free variables of Φ and p appear in Γ . This first judgment is similar to [Aguirre et al. 2017] and the proof rules are thus omitted.

The second judgment $\Gamma; \Phi \vdash e_l : t_l \sim e_r : t_r \mid p$ states that the expressions e_l and e_r , respectively of types t_l and t_r under Γ , satisfy the predicate p , under the assumptions of Φ . To encode this property the predicate p might refer to two special variables r_l and r_r (i.e. not bound in the environment Γ) that respectively refer to the expressions e_l and e_r . For example $\emptyset; \emptyset \vdash 1 : \text{nat} \sim 2 : \text{nat} \mid r_l < r_r$ holds, since $1 < 2$. In general, the relational typing means that $\Gamma; \Phi \vdash p[e_l/r_l][e_r/r_r]$.

Contexts. Contexts are pairs consisting of a typing environment (Γ) that maps variables to their types and a logical environment (Φ) consisting of assertions.

Proof Rules. Figure 8 defines selected proof rules $\Gamma; \Phi \vdash e_l : \text{prM } t_l \sim e_r : \text{prM } t_r \mid p$. The proof rules are given for monadic expressions, in which case the relational assertions are of the form $\diamond_k p$. Informally, these assertions state that 1) the two related expressions satisfy the lifted predicate p and 2) the Kantorovich distance between the two expressions is upper bounded by k .

Rules T-BOT and T-WEAKEN are non-syntax-directed rules that can be applied to any probabilistic expression. Rule T-BOT states that any two probabilistic expressions $e_l \sim e_r$ satisfy the (lifted) true predicate and have expected distance bounded by $+\infty$. Rule T-WEAKEN weakens the lifted predicate $\diamond_k p$ to $\diamond_{k_w} p_w$, provided $p \Rightarrow p_w$ and $k \leq k_w$. These two requirements are established by pure logic.

The rules T-UNIF, T-BERN, and T-CHOICE are used to relate probabilistic primitives. Rule T-UNIF states that the uniform distributions of two non-empty permutations of the same list produce equal values and have zero expected distance. The lifted predicate $\diamond_0 r_l = r_r$ ensures that the two distributions are the same and accordingly their expected distance is zero. We note that more general rules exist, but are omitted here since they are not used in our examples.

The rule T-BERN states that the Bernoulli distribution $\text{bern } x_r$ dominates $\text{bern } x_l$, when $x_l \leq x_r$. The condition is expressed by the pure predicate $r_l \leq r_r$ in the premise, while the conclusion of the rule lifts the same predicate to encode dominance. Finally, it ensures that the distance of the two distributions is bounded by $|x_r - x_l|$.

The rule T-CHOICE relates two choice expressions $\text{choice } x_i \ e_i \ u_i$ (with i being l or r). To do so, it requires that x_l and x_r are equal probabilities, and that the pairs of distributions e_l and e_r and respectively u_l and u_r satisfy the same lifted predicate $\diamond p$ and respectively have distances k_e and k_u . It then ensures that choice will also satisfy the predicate $\diamond p$, while the distance is $x_l \cdot k_e + (1 - x_l) \cdot k_u$.

The rules T-RET and T-BIND are used to relate monadic primitives. The rule T-RET relates $\text{ret } e_l$ with $\text{ret } e_r$. Using pure relational typing, it requires that e_l and e_r satisfy the predicates p (in which the relational variables can freely appear) and their distance is bounded by k . Note that bounding distance in the pure setting is encoded as a logical statement; while a weakening rule can bring the premise in the required syntactic form, potentially with infinite distance. The rule concludes that the return expressions satisfy the lifted p (since the relational variables are not monadic) and their distance is bounded by k . The rule T-BIND relates two expressions $\text{bind } e_i \ f_i \ x_i$, (with i being l or r). In the first premise, it assumes that e_l and e_r satisfy some lifted predicate $\diamond_k q$. In the second premise, the predicate q is assumed in the predicate environment to check the application $f_i \ x_i$ where x_i is the value of the probabilistic argument e_i , i.e. satisfies the lifted predicate $\diamond_{k'} p$, where k' is an affine function of $d(x_l, x_r)$. The bind expressions are then related by $\diamond_{ak+b} p$, where a and b are the coefficients of the affine function.

5.5 Denotational Semantics

Here we define denotational semantics of λ^{RP} by defining the denotations of types and typing environments (§ 5.5.1); expressions (§ 5.5.2); and assertions and predicate environments (§ 5.5.3).

5.5.1 Denotations of Types and Typing Environments. Our denotational semantics only considers discrete distributions, so that expressions have a straightforward set-theoretic interpretation. Definition 5.3 inductively defines the denotations of types. The interesting case is prM that gives a probability distribution, as per definition 5.2.

Definition 5.2 (Discrete Probability Distribution). A probability distribution over a set C is a function $\mu : C \rightarrow [0, 1]$ such that $\sum_{x \in \text{supp } \mu} \mu(x) = 1$. The support of μ is defined as $\text{supp } \mu \doteq \{x \mid x \in C \wedge \mu(x) \neq 0\}$ and needs to be discrete.

We denote the set of discrete distributions over C as $D(C)$.

Definition 5.3 (Denotations of Types). For each type t , $\llbracket t \rrbracket$ is inductively defined as follows:

$$\begin{array}{lll} \llbracket \text{bool} \rrbracket & \doteq & \mathbb{B} \\ \llbracket \text{prM } t \rrbracket & \doteq & D(\llbracket t \rrbracket) \end{array} \quad \begin{array}{lll} \llbracket \text{nat} \rrbracket & \doteq & \mathbb{N} \\ \llbracket \text{real} \rrbracket & \doteq & \mathbb{R}^+ \end{array} \quad \begin{array}{lll} \llbracket \text{list } t \rrbracket & \doteq & \text{list } \llbracket t \rrbracket \\ \llbracket t_x \rightarrow t \rrbracket & \doteq & \llbracket t_x \rrbracket \rightarrow \llbracket t \rrbracket \end{array}$$

Definition 5.4 lifts the denotation of types to typing environment. Concretely, the denotation of a typing environment Γ is a function ρ that maps each declaration $(x : t)$ in Γ to an element in the denotation of t .

Definition 5.4 (Denotation of Type Environments). $\langle \Gamma \rangle \doteq \{\rho \mid \forall (x : t) \in \Gamma. \rho(x) \in \llbracket t \rrbracket\}$

5.5.2 Denotations of Expressions. The denotation of an expression e is defined for a fixed model ρ as $\llbracket e \rrbracket_\rho$ in definition 5.5. The denotations of the pure fragment are standard, where the model is used to define the denotation of a variable as $\rho(x)$ and we skip the verbose definitions for case and let. The probability and return cases use standard distributions, while monadic bind maps to distribution composition, defined in definition 5.6.

Definition 5.5 (Denotations of Expressions). For each expression e and model ρ , $\llbracket e \rrbracket_\rho$ is defined as:

$$\begin{array}{ll} \llbracket \text{unif } e \rrbracket_\rho & \doteq U_{\llbracket e \rrbracket_\rho} \\ \llbracket \text{bern } x \rrbracket_\rho & \doteq B_{\llbracket x \rrbracket_\rho} \\ \llbracket \text{choice } x \ e \ u \rrbracket_\rho & \doteq \llbracket x \rrbracket_\rho \cdot \llbracket e \rrbracket_\rho + (1 - \llbracket x \rrbracket_\rho) \cdot \llbracket u \rrbracket_\rho \\ \llbracket \text{bind } e \ u \rrbracket_\rho & \doteq \text{scomp } \llbracket e \rrbracket_\rho \llbracket u \rrbracket_\rho \\ \llbracket \text{ret } e \rrbracket_\rho & \doteq \delta_{\llbracket e \rrbracket_\rho} \end{array} \quad \begin{array}{ll} \llbracket x \rrbracket_\rho & \doteq \rho(x) \\ \llbracket c \rrbracket_\rho & \doteq c \\ \llbracket \lambda x. e \rrbracket_\rho & \doteq \lambda v. \llbracket e \rrbracket_{\rho[v/x]} \\ \llbracket e \ u \rrbracket_\rho & \doteq \llbracket e \rrbracket_\rho \llbracket u \rrbracket_\rho \end{array}$$

where δ_x represents the Dirac distribution at x , U_{xs} represents the uniform distribution over a non-repeating list xs , and B_p represents the p -biased Bernoulli distribution on $\{0, 1\}$.

Definition 5.6 (Sequential Composition Distribution). Let $\mu \in D(C)$ and $f : C \rightarrow D(C_2)$. Sequential composition distribution $\text{scomp } \mu f$ is defined as:

$$\text{scomp } \mu f(y) \doteq \sum_{x \in C} \mu(x) \cdot f(x)(y)$$

The denotational semantics of expressions is partial, in the sense that some simply typed expressions do not have a denotation. For instance, the interpretation of `unif` is only defined on non-empty lists and the interpretation of `choice` is only defined when x takes values between 0 and 1. It would be possible to make the semantics total by using default values to handle cases not covered by our semantics. However, this would be superfluous because the refinement type system will ensure that all typable expressions have an interpretation.

Our interpretation enjoys a standard soundness theorem.

PROPOSITION 5.7 (SET-THEORETICAL SOUNDNESS). *If $\Gamma \vdash e : t$ and $\rho \in \langle \Gamma \rangle$, then $\llbracket e \rrbracket_\rho \in \llbracket t \rrbracket$.*

5.5.3 Denotations of Assertions and Predicate Environments. Finally, we inductively define denotations of assertions in definition 5.12. Most of the cases are standard except from the case for lifting that relies on Kantorovich couplings (also known as expectation couplings), which we define in definition 5.10 and, in turn, relies on basic definitions of expectation and marginals.

Definition 5.8 (Expectation). For all $\mu \in D(C)$ and $f : C \rightarrow \mathbb{R}$, the expected value $E_{x \leftarrow \mu}[f(x)]$ (or $E_\mu[f]$) of f is a partial function defined as:

$$E_{x \leftarrow \mu}[f(x)] = \sum_{x \in C} \mu(x) \cdot f(x)$$

Definition 5.9 (Marginals). For all $\mu \in D(C_1 \times C_2)$, the first and second marginals of μ are respectively the distributions $\pi_1(\mu) \in D(C_1)$ and $\pi_2(\mu) \in D(C_2)$ defined as:

$$\pi_1(\mu)(x) \doteq \sum_{y \in C_2} \mu(x, y) \quad \pi_2(\mu)(y) \doteq \sum_{x \in C_1} \mu(x, y)$$

Intuitively, the marginals are the “projections” of the couplings, *i.e.* “dependent products” over distributions. Using these, we define Kantorovich coupling, *i.e.* the conditions that render $\diamond_k R$ valid.

Definition 5.10 (Kantorovich coupling). For all $\mu_1 \in D(C_1)$, $\mu_2 \in D(C_2)$, $R \subseteq C_1 \times C_2$, $d : C_1 \times C_2 \rightarrow \mathbb{R}^+$, and $k \in \mathbb{R}^+$, $\models \diamond_k R (\mu_1, \mu_2)$ *iff* there exists $\mu \in D(A_1 \times A_2)$ such that:

$$(1) \pi_1(\mu) = \mu_1 \text{ and } \pi_2(\mu) = \mu_2 \quad (2) \text{supp}(\mu) \subseteq R \quad (3) E_{(x,y) \leftarrow \mu}[d(x, y)] \leq k$$

As usual, $\models p$ means logical validity of p . The first clause corresponds to the standard definition of coupling. The second clause corresponds to the definition of R -coupling and is connected to lifting by the following equivalence: $\diamond R$ *iff* there exists an R -coupling. The last clause is specific to Kantorovich coupling and states that the expected distance with respect to the definition μ is upper bounded by k . Since the Kantorovich distance corresponds to the minimum expected distance over all possible couplings, it follows that k is an upper bound for the Kantorovich distance. This is captured by the following lemma.

LEMMA 5.11. *If $\models \diamond_k R (\mu_1, \mu_2)$, then $\models \diamond R (\mu_1, \mu_2)$ and $\diamond_k \text{true}(\mu_1, \mu_2)$.*

Next, we define the denotation of assertions. We use the Kantorovich distance to interpret distance over distribution types and expectation couplings to interpret lifting, while the rest of the denotations are standard.

Definition 5.12 (Denotation of Assertions). For each assertion p and model ρ , $\llbracket p \rrbracket_\rho$ is defined as:

$$\begin{aligned} \llbracket e \rrbracket_\rho &\doteq \llbracket e \rrbracket_\rho & \llbracket p_1 [\wedge, \vee, \Rightarrow] p_2 \rrbracket_\rho &\doteq \llbracket p_1 \rrbracket_\rho [\wedge, \vee, \Rightarrow] \llbracket p_2 \rrbracket_\rho \\ \llbracket \neg p \rrbracket_\rho &\doteq \neg \llbracket p \rrbracket_\rho & \llbracket \forall x : t. p \rrbracket_\rho &\doteq \forall v : \llbracket t \rrbracket_\rho. \llbracket p \rrbracket_{\rho[v/x]} \\ \llbracket \diamond_k p \ e \ u \rrbracket_\rho &\doteq \diamond_{\llbracket k \rrbracket_\rho} \llbracket p \rrbracket_\rho (\llbracket e \rrbracket_\rho, \llbracket u \rrbracket_\rho) & \llbracket \exists x : t. p \rrbracket_\rho &\doteq \exists v : \llbracket t \rrbracket_\rho. \llbracket p \rrbracket_{\rho[v/x]} \end{aligned}$$

The interpretation of d is defined by induction on types:

$$\begin{aligned} \llbracket d_{\text{bool}} \rrbracket_\rho(x, y) &\doteq \text{if } x = y \text{ then } 0 \text{ else } 1 & \llbracket d_{\text{nat}} \rrbracket_\rho(x, y) &\doteq |x - y| \\ \llbracket d_{\text{real}} \rrbracket_\rho(x, y) &\doteq |x - y| & \llbracket d_{\text{list } t} \rrbracket_\rho(x, y) &\doteq \max_i d_t(x_i, y_i) \\ \llbracket d_{s \rightarrow t} \rrbracket_\rho(x, y) &\doteq \max_{v \in s} d_t(x \ v, y \ v) & \llbracket d_{\text{prM } t} \rrbracket_\rho(x, y) &\doteq \inf_{\mu_{x,y}} E_{\mu_{x,y}}[d] \end{aligned}$$

Here, $\mu_{x,y}$ ranges over couplings of distributions x and y . For lists of equal length, the distance is an index-wise maximum distance between elements, as shown above. Distance between lists of different sizes is infinite.

The denotation of a predicate environment is the conjunction of the denotation of all its predicates.

Definition 5.13 (Denotation of Predicate Environment). $\llbracket \Phi \rrbracket_\rho \doteq \bigwedge_{p \in \Phi} \llbracket p \rrbracket_\rho$

5.6 Soundness

Our proof system is sound with respect to its denotational semantics. That is, for every model of the typing environment that renders the predicate environment valid, the denotation of the predicate, with the special variables r_l and r_r substituted by the typed expressions, is valid.

THEOREM 5.14 (SOUNDNESS). *If $\Gamma; \Phi \vdash e_l : t_l \sim e_r : t_r \mid p$, then for every $\rho \in (\llbracket \Gamma \rrbracket)$ such that $\models \llbracket \Phi \rrbracket_\rho$, we have $\models \llbracket p[\llbracket e_l \rrbracket_\rho / r_l][\llbracket e_r \rrbracket_\rho / r_r] \rrbracket_\rho$.*

In the appendix we prove soundness of the monadic rules, while soundness for the pure fragment follows from Aguirre et al. [2017].

As a corollary of soundness and definition 5.10, we get soundness of the probabilistic fragment. Informally, our judgement relates e_l and e_r when there exists a coupling μ of e_l and e_r such that the predicate p holds for all samples with non-zero probability and the expected distance between samples is less than k .

COROLLARY 5.15 (SOUNDNESS OF PROB. FRAGMENT). *If $\Gamma; \Phi \vdash e_l : \text{prM } t_l \sim e_r : \text{prM } t_r \mid \diamond_k p$, then for every $\rho \in \langle \Gamma \rangle$ such that $\models \llbracket \Phi \rrbracket_\rho$, there exists a $\mu \in D(\llbracket t_l \rrbracket \times \llbracket t_r \rrbracket)$ such that:*

- (1) $\pi_1(\mu) = \llbracket e_l \rrbracket_\rho$ and $\pi_2(\mu) = \llbracket e_r \rrbracket_\rho$;
- (2) $\text{supp}(\mu) \subseteq \{(x_l, x_r) \mid x_l \in \llbracket t_l \rrbracket, x_r \in \llbracket t_r \rrbracket, \models \llbracket p \rrbracket_\rho x_l x_r\}$; and
- (3) $E_{(x,y) \leftarrow \mu}[d(x, y)] \leq \llbracket k \rrbracket_\rho$.

5.7 Derived Rules

By instantiation and combination of the rules that are already present in the system, we can prove other, more specific rules. Such rules are called *derived*. Custom-tailored, they provide shortcuts for verification of common program constructs. Ultimately, extending the system with a variety of derived rules minimizes the proof effort for end users. In this section, we present some of the derived rules of our formalism in fig. 9 and explain their derivation. Only the most practical derived rules have their counterpart in the implementation of our library safe-coupling. Crucially, we don't encode the combined lifting connective $\diamond_k p$. Instead, as discussed in § 3.4 and § 3.5, our assumptions make separate statements about either lifted predicates or distance bounds.

The rule T-BINDRET relates two bind expressions where the second argument is a pure function, composed with `ret` (we use the syntactic sugar for composition $g \cdot f \doteq \lambda x. g (f x)$). Such expressions could be typed by the rule T-BIND followed by T-RET, but this special case permits more convenient reasoning and is used by our case studies. The first premise of the rule is the same as of the T-BIND, but the second premise is now using pure relational typing to bound the distance of f 's result as a function of the distance of its input, i.e. $d(r_l, r_r) \leq m \cdot d(x_l, x_r) + k_u$, so the distance of the conclusion ($m \cdot k_e + k_u$) depends on the distance of the e arguments (k_e).

The rule T-BIND-SIMPL is an instance of the rule T-BIND with the coefficient a set to 0. It assumes that e_l and e_r are related by the lifting of p and that the bodies of the bind map pairs of values that are related by p to pairs of values that are related by $\diamond_k q$. In this case the two monadic binds are also related by $\diamond_k q$.

Rules T-BERN-L and T-BERN-D are obtained from T-BERN by setting the distance to infinity and the predicate to trivial respectively. Such rules are more practical for the implementation from the perspectives of error-reporting and modularity of proofs. In safe-coupling, `bernoulliAxiom` corresponds to T-BERN-L while `bernoulliDist` corresponds to T-BERN-D. Similarly, each rule for probabilistic primitives from fig. 8, as well as T-BIND-SIMPL and T-BINDRET, produce two Liquid Haskell assumptions.

6 RELATED WORK

First Order, Imperative, Probabilistic Languages. There is a large body of work that builds and applies program logics to reason about probabilistic programs. Many of these works are based on first-order imperative languages. Broadly speaking, there exist two main lines of work: the first line of work focuses on non-relational properties and can be traced back to early work by Kozen [1985] and Morgan et al. [1996]. Many of these works focus on establishing sound foundations and have not been implemented in practice. There are however, some noticeable exceptions [Hölzl 2016;

$$\begin{array}{c}
\frac{\Gamma; \Phi \vdash x_l : \text{real} \sim x_r : \text{real} \mid 0 \leq r_l \leq r_r \leq 1}{\Gamma; \Phi \vdash \text{bern } x_l : \text{prM nat} \sim \text{bern } x_r : \text{prM nat} \mid \diamond 0 \leq r_l \leq r_r \leq 1} \text{T-BERN-L} \\
\\
\frac{\Gamma; \Phi \vdash x_l : \text{real} \sim x_r : \text{real} \mid \top}{\Gamma; \Phi \vdash \text{bern } x_l : \text{prM nat} \sim \text{bern } x_r : \text{prM nat} \mid \diamond_{|x_r - x_l|} \top} \text{T-BERN-D} \\
\\
\frac{\begin{array}{c} \Gamma; \Phi \vdash e_l : \text{prM } s \sim e_r : \text{prM } s \mid \diamond q \\ x_r : s, x_l : s, \Gamma; q \ x_l \ x_r, \Phi \vdash f_l \ x_l : \text{prM } t \sim f_r \ x_r : \text{prM } t \mid \diamond_b p \end{array}}{\Gamma; \Phi \vdash \text{bind } e_l \ f_l : \text{prM } t \sim \text{bind } e_r \ f_r : \text{prM } t \mid \diamond_b p} \text{T-BIND-SIMPL} \\
\\
\frac{\begin{array}{c} \Gamma; \Phi \vdash e_l : \text{prM } s_l \sim e_r : \text{prM } s_r \mid \diamond_{k_e} q \\ x_l : s_l, x_r : s_r, \Gamma; q[x_l/r_l][x_r/r_r], \Phi \vdash f_l \ x_l : t_l \sim f_r \ x_r : t_r \mid p \wedge d(r_l, r_r) \leq m \cdot d(x_l, x_r) + k_u \end{array}}{\Gamma; \Phi \vdash \text{bind } e_l \ (\text{ret} \cdot f_l) : \text{prM } t_l \sim \text{bind } e_r \ (\text{ret} \cdot f_r) : \text{prM } t_r \mid \diamond_{m \cdot k_e + k_u} p} \text{T-BINDRET}
\end{array}$$

Fig. 9. Derived rules. We use the syntactic sugar $g \cdot f \doteq \lambda x. g \ (f \ x)$ and $\diamond_k p \doteq \diamond_k (\lambda r_l \ r_r. p) \ r_l \ r_r$.

Hurd 2003], including an active line of work in mechanizing or automating proofs of expected cost of probabilistic programs [Avanzini et al. 2020; Ngo et al. 2018; Tassarotti and Harper 2018].

The second line of work focuses on relational properties; this line of work has been initiated in [Barthe et al. 2009], and its mechanization in the Coq proof assistant or as a self-standing proof assistant [Barthe et al. 2011] have been used to verify formally concrete security of cryptographic constructions. Similar approaches have been developed by the FCF [Petcher and Morrisett 2015] and CryptHOL [Basin et al. 2017]. These logics have been further extended to reason about differential privacy [Barthe et al. 2012].

Our work is most closely related to work on expected sensitivity and in particular [Aguirre et al. 2021b; Barthe et al. 2018]. Barthe et al. [2018] introduce the notion of expectation coupling used in this paper and show how expectation couplings can be used to prove expected sensitivity. In addition, they provide a relational program logic to prove the existence of expectation couplings between two probabilistic imperative programs. The program logic manipulates assertions of the form p, d where p is a relational boolean-valued assertion and d is a relational quantitative assertion. The proof system departs from classic program logics by being highly not syntax-directed; for instance, the rule for sequential composition combines the classic rule with a case analysis. This makes application of the rule rather complex. Moreover the proof system is not implemented. The program logic is used to prove stability of SGD; however the proof is on paper and not mechanized. Our work can be construed as an attempt to raise (a streamlined version of) the program logic from [Barthe et al. 2018] to a higher-order setting. As pointed out in [Aguirre et al. 2021a], lifting quantitative relational logics to the higher-order setting is far from straightforward; in the terminology of Aguirre et al. [2021a], it seems difficult to define a $[0, +\infty]$ relational lifting to instantiate their generic program logic. We forego this issue by restricting our language so that we can give a direct, set-theoretic, semantics of programs.

Aguirre et al. [2021b] develop a relational weakest pre-expectation calculus for reasoning about expected sensitivity of a probabilistic imperative language. Their pre-expectation calculus can be construed as a relational variant of pGCL [Morgan et al. 1996]. The main novelty of their calculus is the rule for random sampling, which takes the minimal relational pre-expectation over all possible couplings of the two sampled distributions. They show that their calculus subsumes a fragment of the logic from [Barthe et al. 2018]. Their calculus is used to verify our two main examples. Interestingly, their proof of convergence of TD(0) is significantly different from ours. Our alternative proof is

based on vanilla couplings and establishes a stronger statement: we prove that the distance between the two outputs is upper bounded for all elements of the coupling, whereas they prove that the expected distance between the two outputs is upper bounded for the same coupling. Thus, our statement entails theirs by basic properties of expectations. Lifting pre-expectation calculus to a higher-order setting is challenging, even in the non-relational setting. Indeed, [Avanzini et al. \[2021\]](#) develop a two-step approach for reasoning formally about expectations, and in particular expected cost, of higher-order probabilistic programs. In the first step, probabilistic programs are transformed into deterministic programs using a variant of the well-known continuation-passing style (CPS) translation. They show that the translation is faithful: for instance, the expected cost of a program can be recovered exactly from its translation, by applying the translation to the continuation $\lambda\kappa.0$. In the second step, they use a refinement type system to prove upper bounds of (real-valued) expressions. The refinement type system uses the notion of admissible predicate to ensure that the classic rule for fixpoints remains sound in presence of non-terminating computations. Extending [\[Avanzini et al. 2021\]](#) to a relational setting is an interesting direction for future work.

There are further works that consider expected sensitivity of probabilistic programs. For instance, [Wang et al. \[2020\]](#) develop an alternative method based on martingale theory for proving expected sensitivity of probabilistic programs. Their analysis covers a rich set of termination behaviors, in particular it covers programs with probabilistic loop guards. One drawback of their approach is that it proves a weaker property, namely expected sensitivity for some finite constant. Their analysis is implemented and evaluated on several examples, but neither SGD nor TD. Independently, [Huang et al. \[2018\]](#) propose another automated approach based on symbolic computation to analyze the expected sensitivity of programs. They also do not analyze SGD or TD.

Relational Reasoning for Higher Order Languages. [Nanevski et al. \[2011\]](#) were among the first to explore relational reasoning for higher-order languages. Their work defines Relational Hoare Type Theory (RHTT), a powerful program logic for proving relational properties of stateful higher-order programs. RHTT is implemented in the Coq proof assistant and is used to verify intricate information flow properties. The main difference with our work is that RHTT operates on a shallow embedding of programs and does not support probabilistic reasoning. BiRelCost [\[Çiçek et al. 2019\]](#) is a bidirectional type checker that automatically performs relational and unary cost analysis requiring minimal user annotations. However, the checker is incomplete and relies on example-driven heuristics. It is unlikely that the used heuristics would suffice to prove our case studies. [Handley et al. \[2019\]](#), like us, address this incompleteness by encoding the relational rules in Liquid Haskell and requiring explicit, user-provided, extrinsic proofs; sacrificing automation in the name of expressiveness and predictability. Our work applies the technique of [Handley et al. \[2019\]](#) to reason about probabilistic programs and quantitative specifications.

Verification of Higher Order Probabilistic Programs. There have been two lines of work that verify relational properties on executable probabilistic programs in F^* [\[Swamy et al. 2016\]](#). First, rF^* [\[Barthe et al. 2014\]](#) is an extension of F^* that supports relational reasoning via relational refinement types. As with F^* , type checking generates SMT queries that are discharged by Z3. Although rF^* supports probabilistic programs, reasoning is constrained by syntax-directed typing, for example the `binsDist` of § 2.5 cannot be type-checked. We overcome this limitation by supporting a richer set of axioms and extrinsic proofs. Second, [Grimm et al. \[2018\]](#) present an alternative approach to reason about relational properties in F^* that, similar to our work, uses extrinsic proofs to reason about programs. Their approach supports probabilistic reasoning and is used to prove probabilistic non-interference of Shannon’s classic cryptographic one-time pad. However, their support for probabilistic reasoning is limited and quantitative specifications are not supported.

Proof Systems of Higher Order Probabilistic Programs. Aguirre et al. [2021a] present several proof systems for higher-order stateful probabilistic programs. One key novelty of their proof systems is the support of interactive adversarial computations. Such rules can be used to prove that programs verify relational properties for all possible (well-typed) adversarial computations. Unfortunately, these rules are only proved sound for boolean relational properties and it is an open problem whether these rules remain sound for quantitative properties modeled using Kantorovich lifting. Their approach builds on Relational Higher-Order Logic (RHOL) [Aguirre et al. 2017], which achieves full expressiveness and bypasses limitations of syntax-directed reasoning via an embedding into HOL. However, there is no implementation of RHOL and of its extensions.

Differential privacy. Differential privacy is a mathematical notion of privacy that can be understood as a form of probabilistic sensitivity *w.r.t.* some pseudo-distance induced by a specific f -divergence. As such, differential privacy is directly related to our work.

Fuzz [Reed and Pierce 2010] was the first type-based approach to verify distance of higher order programs, in the domain of differential privacy. Since introduced, Fuzz has been extended in various ways. DFuzz [Gaboardi et al. 2013] introduces recursion; Adaptive Fuzz [Winograd-Cort et al. 2017] supports dynamic data analysis; Fuzzi [Zhang et al. 2019] extend Fuzz with APRHL [Barthe et al. 2012] to prove trusted primitives (like Laplace); Duet [Near et al. 2019] extends Fuzz with support for advanced variants of differential privacy via a dual type system; HOARe2 [Barthe et al. 2015] provides a relational refinement system that embeds Fuzz; and Bunched Fuzz [June Wunder and Arthur Azevedo de Amorim and Patrick Baillot and Marco Gaboardi 2022] extends the system with bunches to, like us, reason about distances of probability distributions. Like our system, most of these approaches use typing rules to trace distance and have some support of probabilistic programs. However, they lack the flexibility to reason about expected sensitivity for most advanced examples such as those considered here.

In a work most closely related to ours, DPella [Lobo-Vesga et al. 2021] and Solo [Abuah et al. 2021] use Haskell's dependent types to encode differential privacy. Both these systems are similar to ours since they do verify executable Haskell code. The critical difference is that they use Haskell's dependent types while we use the refinement type extension of Liquid Haskell. One benefit of our approach is that we can delegate arithmetic reasoning to SMT. At a more general level, the two approaches are incomparable: we are able to reason finely about expected sensitivity, whereas they can reason about differential privacy using standard composition theorems, and about accuracy using a fine-grained approach that exploits a slick combination of information flow typing and concentration inequalities.

7 CONCLUSION

We presented `safe-coupling`, a library that allows reasoning about relational properties of probabilistic, executable, Haskell programs. `safe-coupling` is developed on top of Liquid Haskell and as such enjoys the predictable verification of a mature, unary refinement type checker. To justify the assumptions of our library we formalized the core calculus λ^{RP} that captures these assumptions and proved it sound. Finally, we used `safe-coupling` to formally prove convergence of TD(0) and algorithmic stability of SGD, both classic machine learning algorithms from the literature.

ACKNOWLEDGMENTS

We thank Alejandro Aguirre and Jiajia Song for the useful discussions. This work is funded by the Horizon Europe ERC Starting Grant CRETE (GA: 101039196), the US Office of Naval Research HACKCRYPT (Ref. N00014-19-1-2292), the Atracción de Talento grant (Ref. 2019-T2/TIC-13455), and the Juan de la Cierva grant (IJC2019-041599-I).

REFERENCES

- Chike Abuah, David Darais, and Joseph P. Near. 2021. Solo: Enforcing Differential Privacy Without Fancy Types. In *CoRR*. <https://arxiv.org/abs/2105.01632>
- Alejandro Aguirre, Gilles Barthe, Marco Gaboardi, Deepak Garg, Shin-ya Katsumata, and Tetsuya Sato. 2021a. Higher-Order Probabilistic Adversarial Computations: Categorical Semantics and Program Logics. In *ICFP*. <https://doi.org/10.1145/3473598>
- Alejandro Aguirre, Gilles Barthe, Marco Gaboardi, Deepak Garg, and Pierre-Yves Strub. 2017. A Relational Logic for Higher-Order Programs. In *ICFP*. <https://doi.org/10.1145/3110265>
- Alejandro Aguirre, Gilles Barthe, Justin Hsu, Benjamin Lucien Kaminski, Joost-Pieter Katoen, and Christoph Matheja. 2021b. A Pre-Expectation Calculus for Probabilistic Sensitivity. In *POPL*. <https://doi.org/10.1145/3434333>
- Martin Avanzini, Gilles Barthe, and Ugo Dal Lago. 2021. On Continuation-Passing Transformations and Expected Cost Analysis. In *ICFP*. <https://doi.org/10.1145/3473592>
- Martin Avanzini, Georg Moser, and Michael Schaper. 2020. A Modular Cost Analysis for Probabilistic Programs. In *OOPSLA*. <https://doi.org/10.1145/3428240>
- Gilles Barthe, Thomas Espitau, Benjamin Grégoire, Justin Hsu, and Pierre-Yves Strub. 2018. Proving Expected Sensitivity of Probabilistic Programs. In *POPL*. <https://doi.org/10.1145/3158145>
- Gilles Barthe, Cédric Fournet, Benjamin Grégoire, Pierre-Yves Strub, Nikhil Swamy, and Santiago Zanella Béguelin. 2014. Probabilistic Relational Verification for Cryptographic Implementations. In *POPL*. <https://doi.org/10.1145/2535838.2535847>
- Gilles Barthe, Marco Gaboardi, Emilio Jesús Gallego Arias, Justin Hsu, Aaron Roth, and Pierre-Yves Strub. 2015. Higher-Order Approximate Relational Refinement Types for Mechanism Design and Differential Privacy. In *POPL*. <https://doi.org/10.1145/2676726.2677000>
- Gilles Barthe, Benjamin Grégoire, and Santiago Zanella Béguelin. 2009. Formal Certification of Code-Based Cryptographic Proofs. In *POPL*. <https://doi.org/10.1145/1480881.1480894>
- Gilles Barthe, Benjamin Grégoire, Sylvain Heraud, and Santiago Zanella Béguelin. 2011. Computer-Aided Security Proofs for the Working Cryptographer. In *CRYPTO*. https://doi.org/10.1007/978-3-642-22792-9_5
- Gilles Barthe and Justin Hsu. 2020. *Probabilistic Couplings from Program Logics*. Cambridge University Press. <https://doi.org/10.1017/9781108770750.006>
- Gilles Barthe, Boris Köpf, Federico Olmedo, and Santiago Zanella Béguelin. 2012. Probabilistic Relational Reasoning for Differential Privacy. In *POPL*. <https://doi.org/10.1145/2103656.2103670>
- David A. Basin, Andreas Lochbihler, and S. Reza Sefidgar. 2017. CryptHOL: Game-based Proofs in Higher-order Logic. In *Journal of Cryptology*. <https://doi.org/10.1007/s00145-019-09341-z>
- Olivier Bousquet and André Elisseeff. 2002. Stability and Generalization. In *Journal of Machine Learning Research*. <https://doi.org/10.1162/153244302760200704>
- Ezgi Çiçek, Weihao Qu, Gilles Barthe, Marco Gaboardi, and Deepak Garg. 2019. Bidirectional Type Checking for Relational Properties. In *PLDI*. <https://doi.org/10.1145/3314221.3314603>
- Yuxin Deng. 2015. *Semantics of Probabilistic Processes: An Operational Approach*. Springer. <https://doi.org/10.1007/978-3-662-45198-4>
- Marco Gaboardi, Andreas Haeberlen, Justin Hsu, Arjun Narayan, and Benjamin C. Pierce. 2013. Linear Dependent Types for Differential Privacy. In *POPL*. <https://doi.org/10.1145/2429069.2429113>
- Niklas Grimm, Kenji Maillard, Cédric Fournet, Catalin Hritcu, Matteo Maffei, Jonathan Protzenko, Tahina Ramananandro, Aseem Rastogi, Nikhil Swamy, and Santiago Zanella Béguelin. 2018. A Monadic Framework for Relational Verification: Applied to Information Security, Program Equivalence, and Optimizations. In *CPP*. <https://doi.org/10.1145/3167090>
- Jad Hamza, Nicolas Voirol, and Viktor Kunčák. 2019. System FR: Formalized Foundations for the Stainless Verifier. In *OOPSLA*. <https://doi.org/10.1145/3360592>
- Martin A. T. Handley, Niki Vazou, and Graham Hutton. 2019. Liquidate Your Assets: Reasoning about Resource Usage in Liquid Haskell. In *POPL*. <https://doi.org/10.1145/3371092>
- Moritz Hardt, Ben Recht, and Yoram Singer. 2016. Train Faster, Generalize Better: Stability of Stochastic Gradient Descent. In *ICML*. <https://dl.acm.org/doi/10.5555/3045390.3045520>
- Johannes Hölzl. 2016. Formalising Semantics for Expected Running Time of Probabilistic Programs. In *ITP*. https://doi.org/10.1007/978-3-319-43144-4_30
- Zixin Huang, Zhenbang Wang, and Sasa Misailovic. 2018. PSense: Automatic Sensitivity Analysis for Probabilistic Programs. In *ATVA*. https://doi.org/10.1007/978-3-030-01090-4_23
- Joe Hurd. 2003. Verification of the Miller-Rabin probabilistic primality test. In *The Journal of Logic and Algebraic Programming*. [https://doi.org/10.1016/S1567-8326\(02\)00065-6](https://doi.org/10.1016/S1567-8326(02)00065-6)
- june wunder and Arthur Azevedo de Amorim and Patrick Baillot and Marco Gaboardi. 2022. Bunched Fuzz: Sensitivity for Vector Metrics. In *CoRR*. <https://doi.org/10.48550/arXiv.2202.01901>

- Dexter Kozen. 1985. A Probabilistic PDL. In *Journal of Computer and System Sciences*. [https://doi.org/10.1016/0022-0000\(85\)90012-1](https://doi.org/10.1016/0022-0000(85)90012-1)
- Nico Lehmann, Rose Kunkel, Jordan Brown, Jean Yang, Niki Vazou, Nadia Polikarpova, Deian Stefan, and Ranjit Jhala. 2021. STORM: Refinement Types for Secure Web Applications. In *OSDI*. <https://www.usenix.org/conference/osdi21/presentation/lehmann>
- Torgny Lindvall. 2002. *Lectures on the Coupling Method*. <https://doi.org/10.1137/1035121>
- Elisabet Lobo-Vesga, Alejandro Russo, and Marco Gaboardi. 2021. A Programming Language for Data Privacy with Accuracy Estimations. In *Transactions on Programming Languages and Systems*. <https://doi.org/10.1145/3452096>
- Kenji Maillard, Catalin Hritcu, Exequiel Rivas, and Antoine Van Muylder. 2020. The Next 700 Relational Program Logics. In *POPL*. <https://doi.org/10.1145/3371072>
- Carroll Morgan, Annabelle McIver, and Karen Seidel. 1996. Probabilistic Predicate Transformers. In *TOPLAS*. <https://doi.org/10.1145/229542.229547>
- Aleksandar Nanevski, Anindya Banerjee, and Deepak Garg. 2011. Verification of Information Flow and Access Control Policies with Dependent Types. In *S&P*. <https://doi.org/10.1109/SP.2011.12>
- Joseph P. Near, David Darais, Chike Abuah, Tim Stevens, Pranav Gaddamadugu, Lun Wang, Neel Somani, Mu Zhang, Nikhil Sharma, Alex Shan, and Dawn Song. 2019. Duet: An Expressive Higher-Order Language and Linear Type System for Statically Enforcing Differential Privacy. (2019). <https://doi.org/10.1145/3360598>
- Van Chan Ngo, Quentin Carbonneaux, and Jan Hoffmann. 2018. Bounded Expectations: Resource Analysis for Probabilistic Programs. In *PLDI*. <https://doi.org/10.1145/3192366.3192394>
- Adam Petcher and Greg Morrisett. 2015. The Foundational Cryptography Framework. In *POST*. https://doi.org/10.1007/978-3-662-46666-7_4
- Norman Ramsey and Avi Pfeffer. 2002. Stochastic Lambda Calculus and Monads of Probability Distributions. In *POPL*. <https://doi.org/10.1145/503272.503288>
- Jason Reed and Benjamin C. Pierce. 2010. Distance Makes the Types Grow Stronger: A Calculus for Differential Privacy. In *ICFP*. <https://doi.org/10.1145/1863543.1863568>
- Adam Scibior, Zoubin Ghahramani, and Andrew D. Gordon. 2015. Practical Probabilistic Programming with Monads. In *Haskell*. <https://doi.org/10.1145/2887747.2804317>
- Nikhil Swamy, Catalin Hritcu, Chantal Keller, Aseem Rastogi, Antoine Delignat-Lavaud, Simon Forest, Karthikeyan Bhargavan, Cédric Fournet, Pierre-Yves Strub, Markulf Kohlweiss, Jean-Karim Zinzindohoué, and Santiago Zanella-Béguelin. 2016. Dependent Types and Multi-Monadic Effects in F*. In *POPL*. <https://doi.org/10.1145/2914770.2837655>
- Joseph Tassarotti and Robert Harper. 2018. Verified Tail Bounds for Randomized Programs. In *ITP*. https://doi.org/10.1007/978-3-319-94821-8_33
- Hermann Thorisson. 2000. *Coupling, Stationarity, and Regeneration*. Springer. <https://notendur.hi.is/hermann/iid/csr/>
- Elizaveta Vasilenko and Niki Vazou. 2022. Safe Couplings: Coupled Refinement Types. Zenodo. <https://doi.org/10.5281/zenodo.6710298>
- Niki Vazou, Joachim Breitner, Rose Kunkel, David Van Horn, and Graham Hutton. 2018. Theorem Proving for All: Equational Reasoning in Liquid Haskell (Functional Pearl). In *Haskell*. <https://doi.org/10.1145/3242744.3242756>
- Niki Vazou, Eric L. Seidel, and Ranjit Jhala. 2014a. LiquidHaskell: Experience with Refinement Types in the Real World. In *Haskell*. <https://doi.org/10.1145/2775050.2633366>
- Niki Vazou, Eric L. Seidel, Ranjit Jhala, Dimitrios Vytiniotis, and Simon Peyton-Jones. 2014b. Refinement Types for Haskell. In *ICFP*. <https://doi.org/10.1145/2692915.2628161>
- Niki Vazou, Anish Tondwalkar, Vikraman Choudhury, Ryan G. Scott, Ryan R. Newton, Philip Wadler, and Ranjit Jhala. 2017. Refinement Reflection: Complete Verification with SMT. In *POPL*. <https://doi.org/10.1145/3158141>
- Cédric Villani. 2009. *Optimal Transport, old and new*. Springer. <https://link.springer.com/book/10.1007/978-3-540-71050-9>
- Peixin Wang, Hongfei Fu, Krishnendu Chatterjee, Yuxin Deng, and Ming Xu. 2020. Proving Expected Sensitivity of Probabilistic Programs with Randomized Variable-Dependent Termination Time. In *POPL*. <https://doi.org/10.1145/3371093>
- Daniel Winograd-Cort, Andreas Haeberlen, Aaron Roth, and Benjamin C. Pierce. 2017. A Framework for Adaptive Differential Privacy. In *ICFP*. <https://doi.org/10.1145/3110254>
- Hengchu Zhang, Edo Roth, Andreas Haeberlen, Benjamin C. Pierce, and Aaron Roth. 2019. Fuzzi: A Three-Level Logic for Differential Privacy. In *ICFP*. <https://doi.org/10.1145/3341697>